

DEEP LEARNING: IMAGE CLASSIFICATION

Chih-Chung Hsu (許志仲)

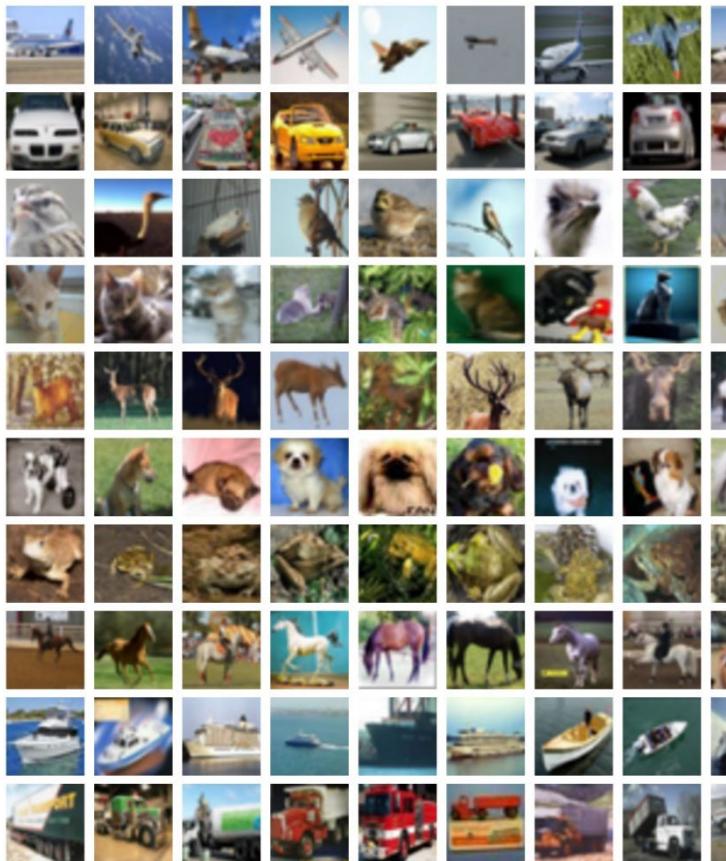
Institute of Data Science

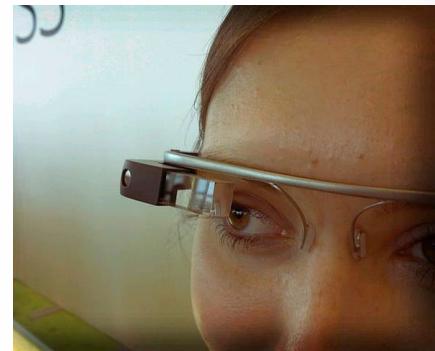
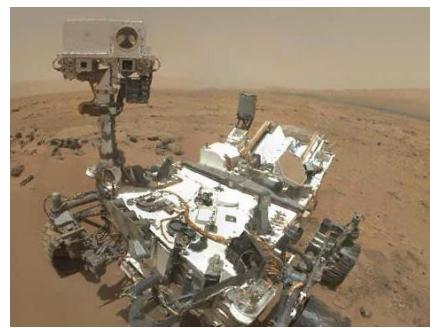
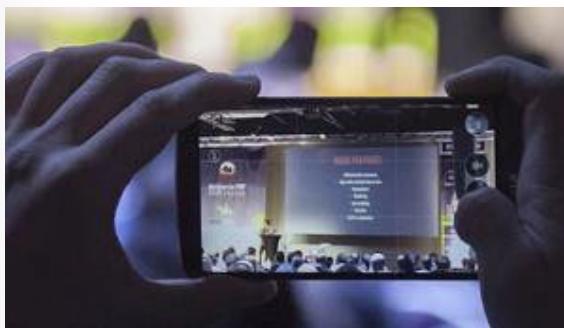
National Cheng Kung University

<https://cchsu.info>



file





Top row, left to right:

[Image](#) by [Roger H Goun](#) is licensed under [CC BY 2.0](#)

[Image](#) is [CC0 1.0](#) public domain

[Image](#) is [CC0 1.0](#) public domain

[Image](#) is [CC0 1.0](#) public domain

Middle row, left to right

[Image](#) by [BGPHP Conference](#) is licensed under [CC BY 2.0](#); changes made

[Image](#) is [CC0 1.0](#) public domain

[Image](#) by [NASA](#) is licensed under [CC BY 2.0](#)

[Image](#) is [CC0 1.0](#) public domain

Bottom row, left to right [Image](#)

is [CC0 1.0](#) public domain

[Image](#) by [Derek Keats](#) is licensed under [CC BY 2.0](#); changes made

[Image](#) is public domain

[Image](#) is licensed under [CC-BY 2.0](#); changes made

Biology

Psychology

Physics

**Computer
Science**

Engineering

Mathematics

**Computer
Vision**

Neuroscience

optics

Cognitive
sciences

graphics, algorithms,
theory,...

systems,
architecture, ...

Image
processing

Speech, NLP

Robotics

Information retrieval

Machine learning

Evolution's Big Bang



[This image is licensed under CC-BY 2.5](#)



[This image is licensed under CC-BY 2.5](#)



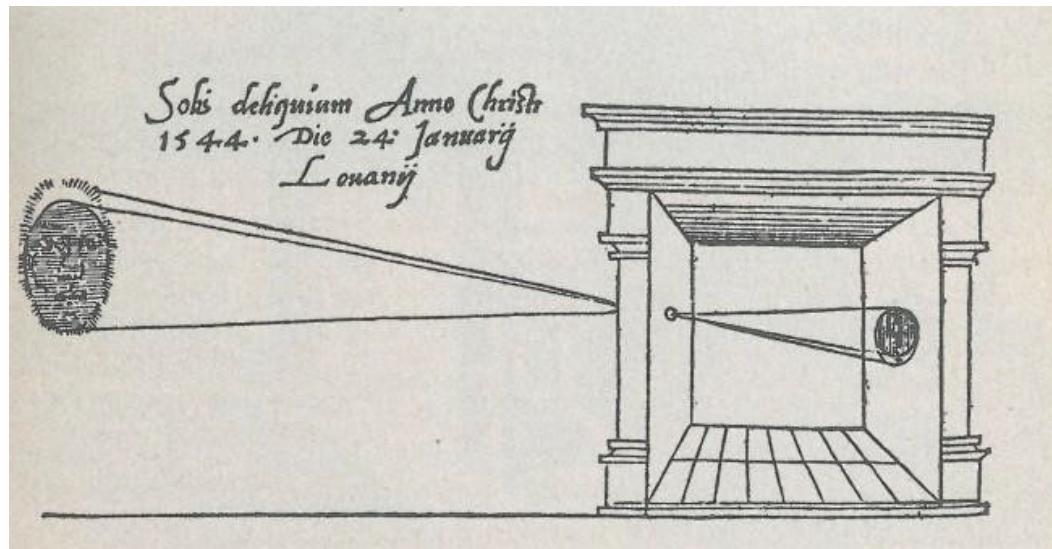
[This image is licensed under CC-BY 3.0](#)

543 million years, B.C.

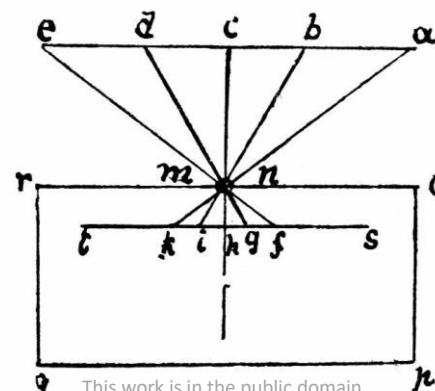
ACVLab@NCKU

Camera Obscura

Gemma Frisius, 1545



[This work is in the public domain](#)

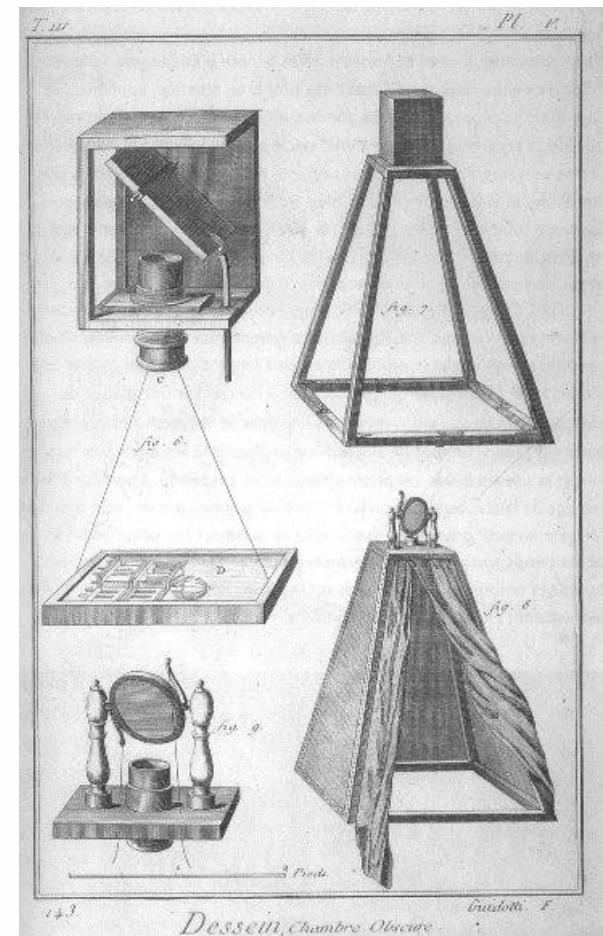


[This work is in the public domain](#)

2023/2/22

Leonardo da Vinci,
16th Century AD

Encyclopedie, 18th Century



[This work is in the public domain](#)

MASSACHUSETTS INSTITUTE OF TECHNOLOGY
PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.

July 7, 1966

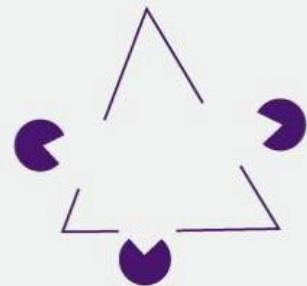
THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

Copyrighted Material

VISION



David Marr

FOREWORD BY
Shimon Ullman

AFTERWORD BY
Tomaso Poggio

Copyrighted Material

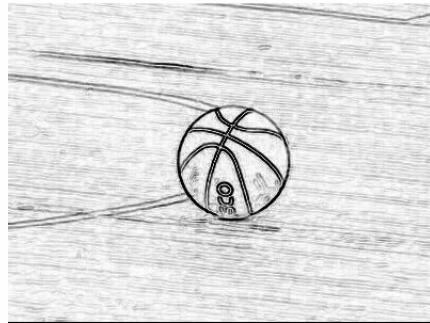
David Marr, 1970s

Input image

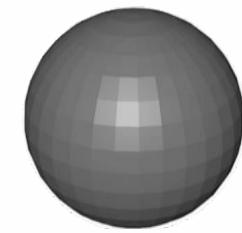
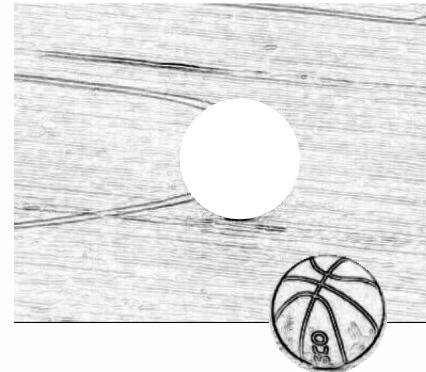


[This image is CC0 1.0 public domain](#)

Edge image

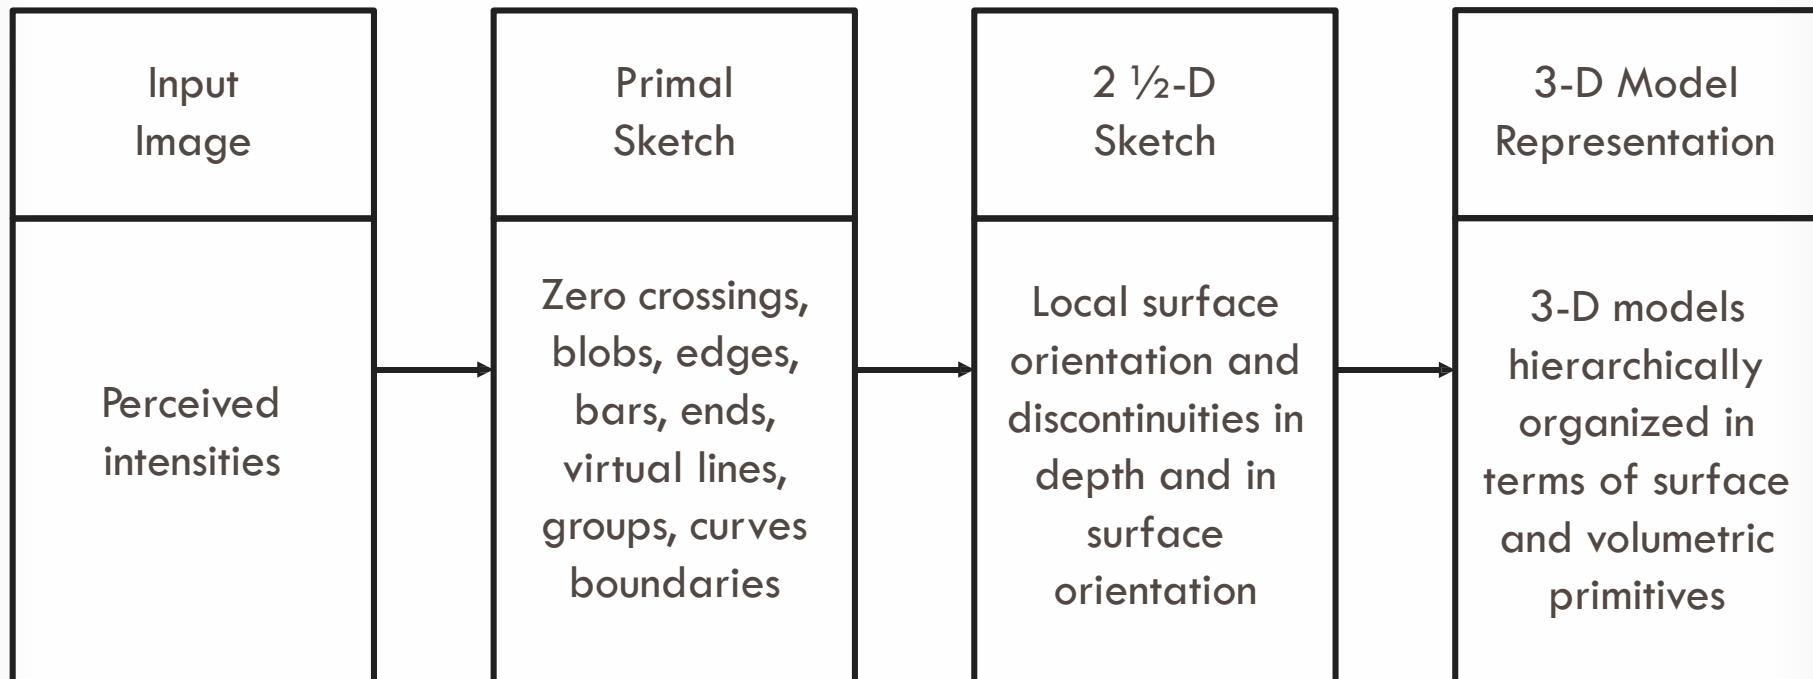


3-D model



[This image is CC0 1.0 public domain](#)

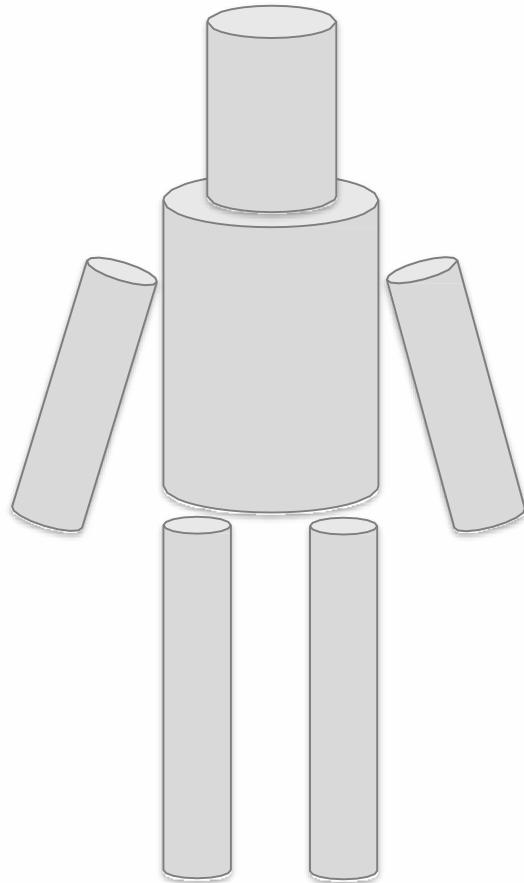
2 1/2-D sketch



Stages of Visual Representation, David Marr, 1970s

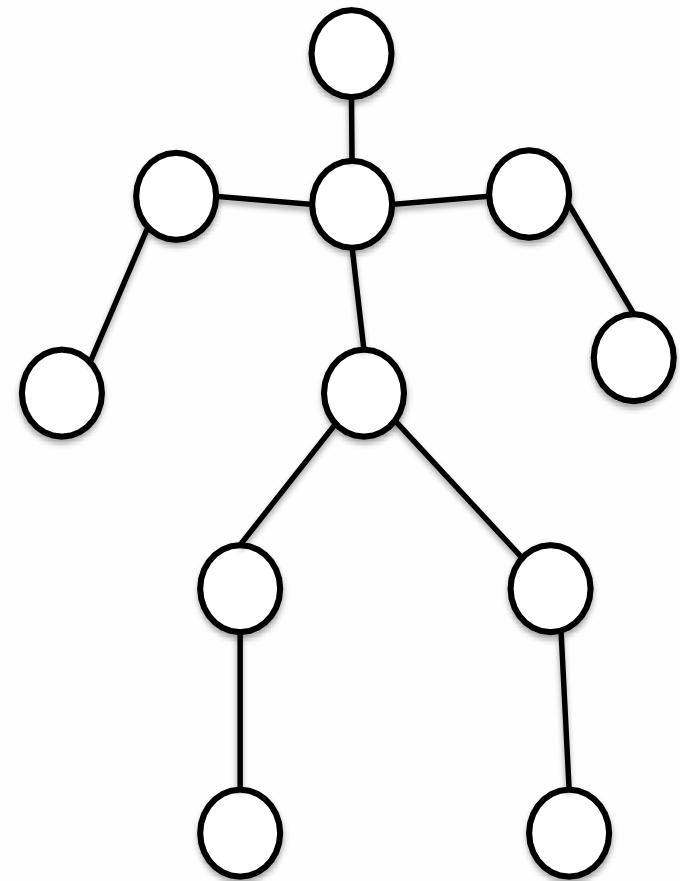
- Generalized Cylinder

Brooks & Binford, 1979



- Pictorial Structure

Fischler and Elschlager, 1973



Edge map to form the objects



David Lowe, 1987

Normalized Cut (Shi & Malik, 1997)

[Image is CC BY3.0](#)



[Image is public domain](#)



[Image is CC-BY2.0; changes made](#)



Face Detection, Viola & Jones,
2001

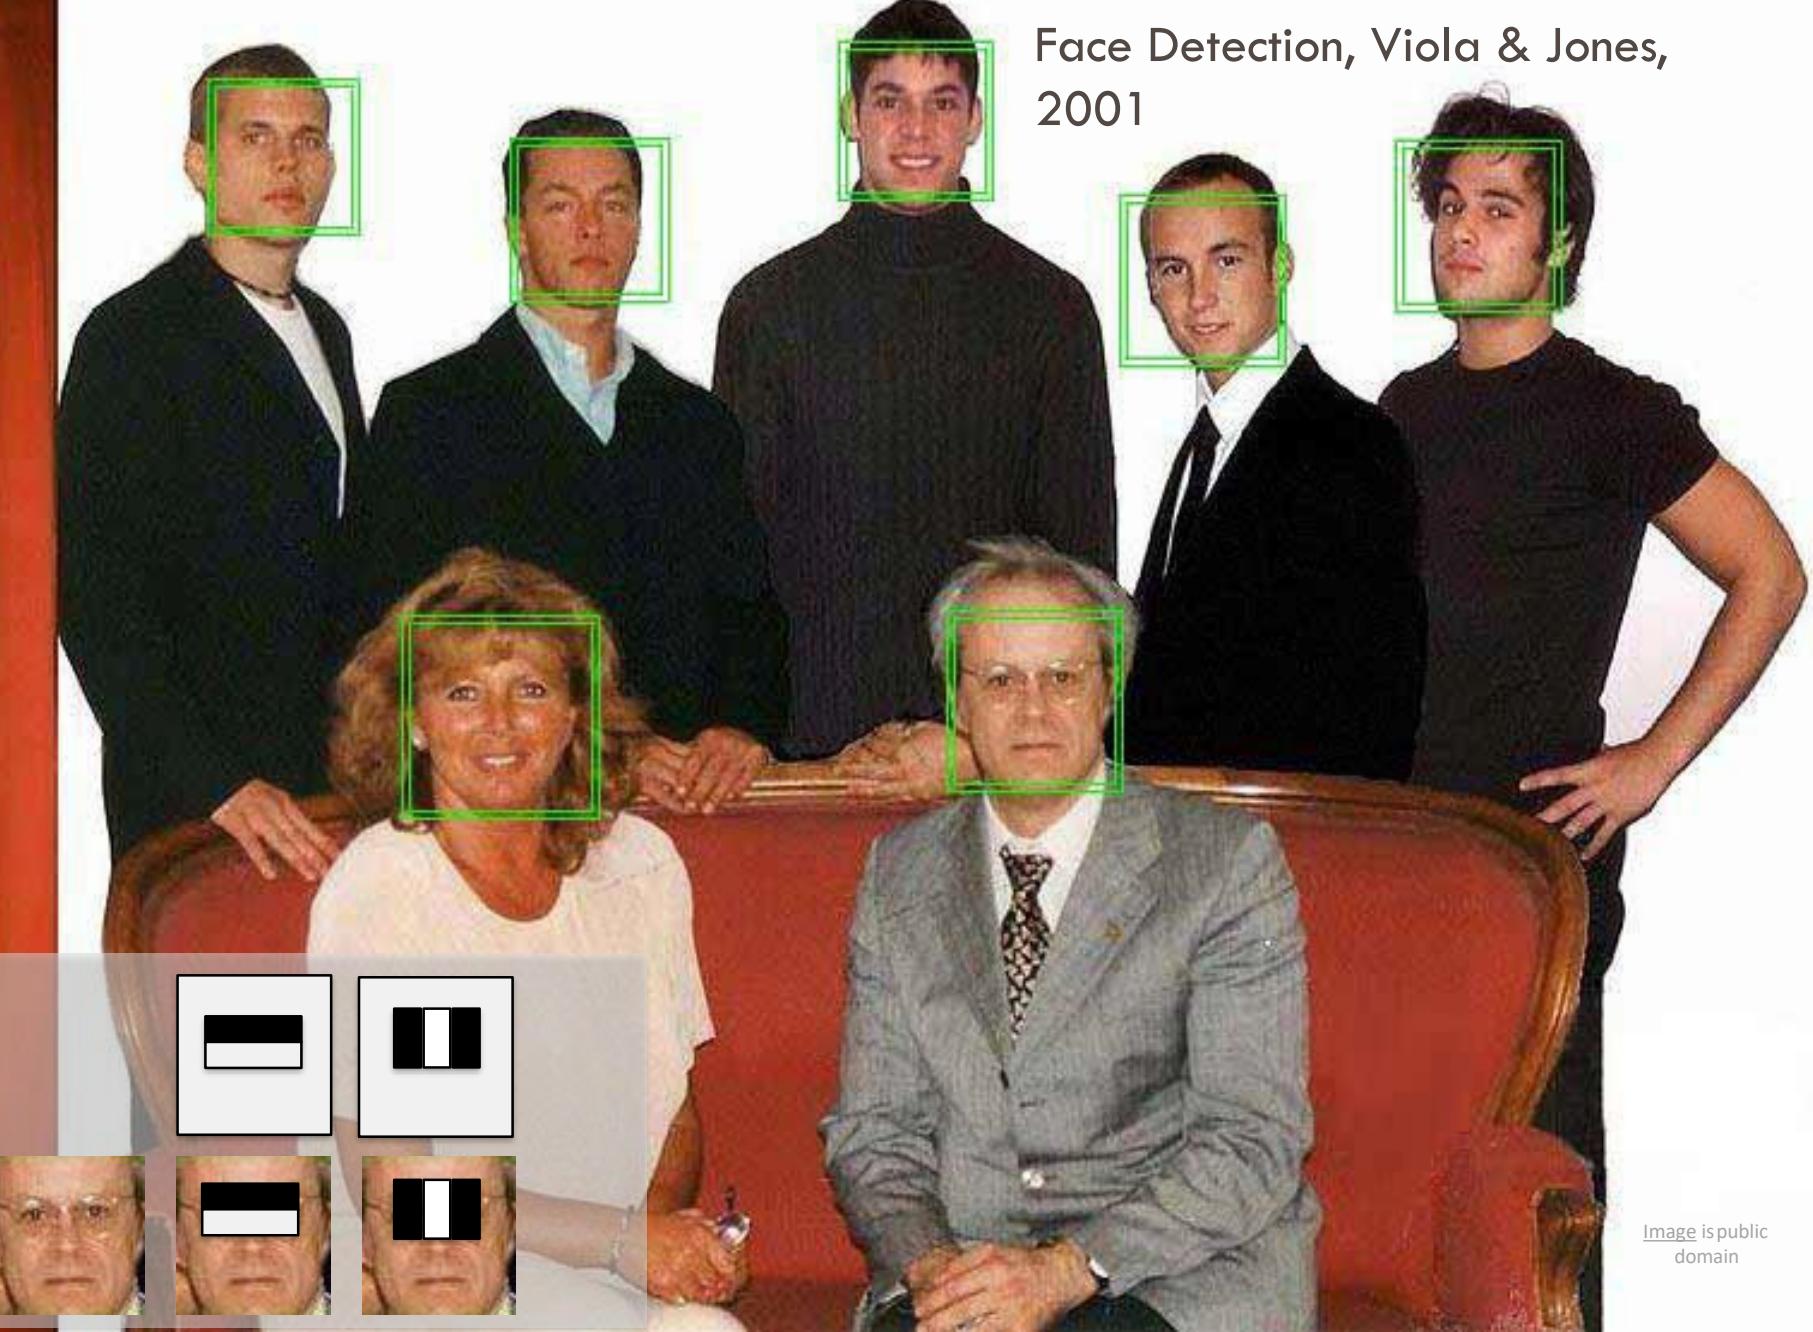


Image is public
domain

Local Feature and Matching



[Image](#) is public domain



[Image](#) is public domain

“SIFT” & Object Recognition, David Lowe, 1999

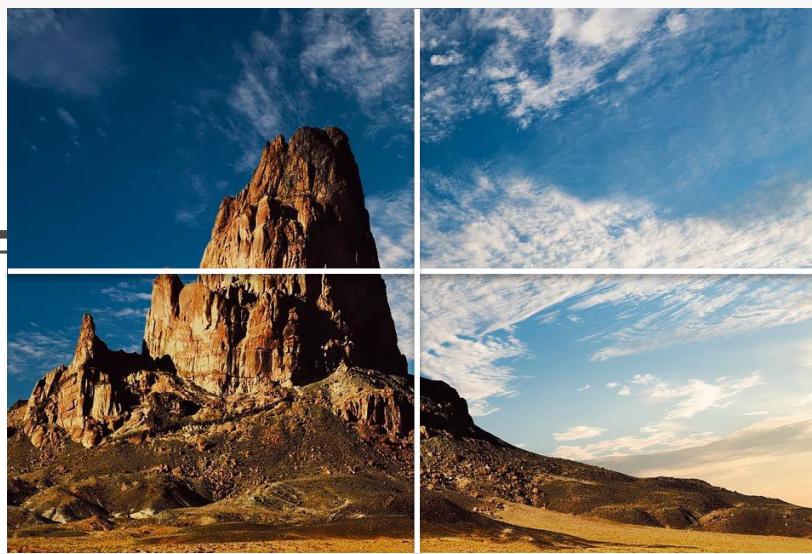
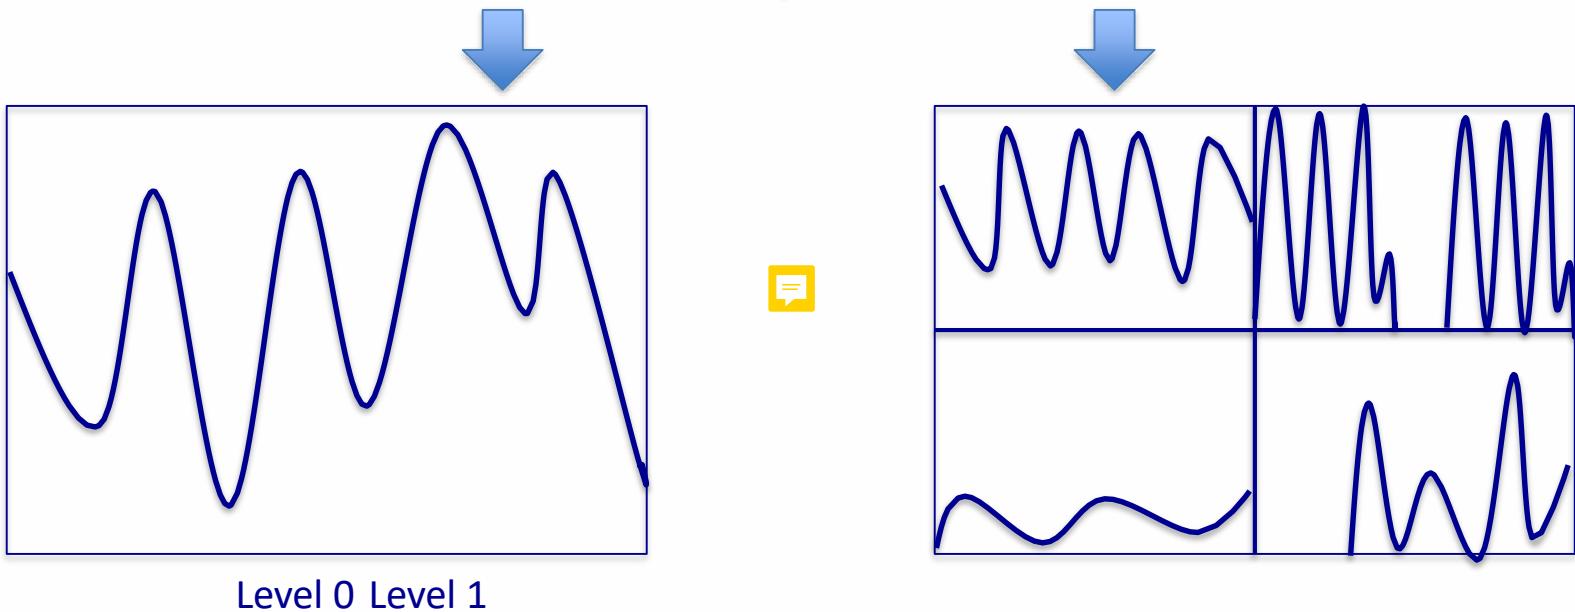
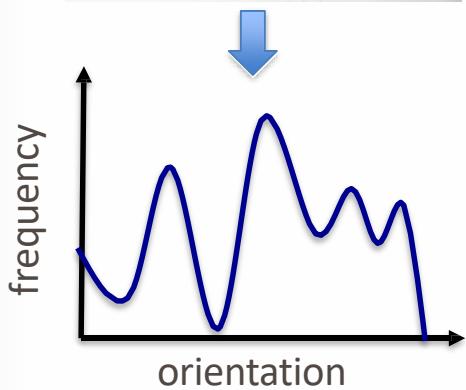


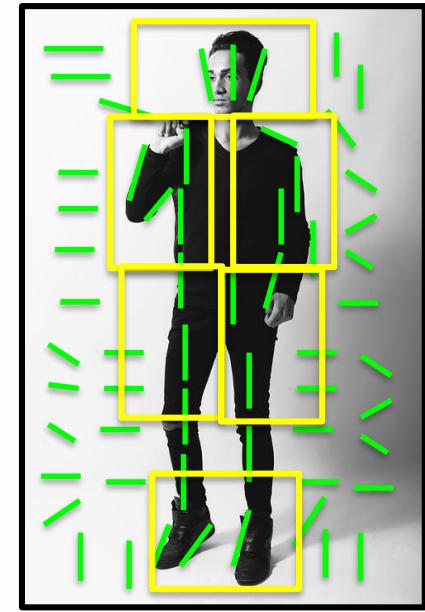
Image is CC0 1.0 public domain



Spatial Pyramid Matching, Lazebnik, Schmid & Ponce, 2006



Histogram of Gradients (HoG)
Dalal & Triggs, 2005



Deformable Part Model
Felzenswalb, McAllester, Ramanan, 2009

PASCAL Visual Object Challenge

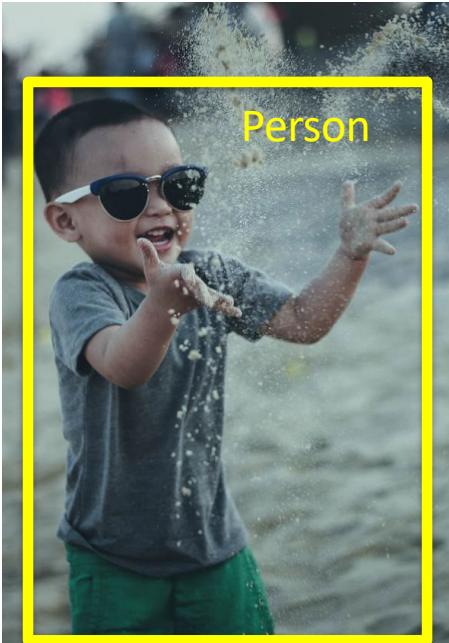
(20 object categories)

[Everingham et al. 2006-2012]

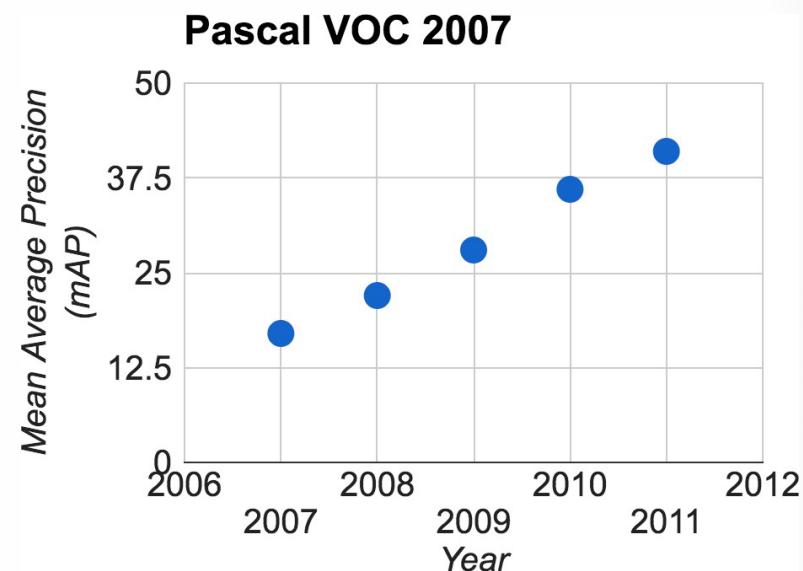
[Image](#) is [CC0 1.0](#) public domain



[Image](#) is [CC0 1.0](#) public domain



[Image](#) is [CC0 1.0](#) public domain





www.image-net.org

22K categories and **15M** images

- Animals
 - Bird
 - Fish
 - Mammal
 - Invertebrate
- Plants
 - Tree
 - Flower
- Food
- Materials
- Structures
- Artifact
 - Tools
 - Appliances
 - Structures
- Person
- Scenes
 - Indoor
 - Geological Formations
- Sport Activities

Deng, Dong, Socher, Li, Li, & Fei-Fei, 2009

The Image Classification Challenge:
1,000 object classes
1,431,167 images



Output:
Scale
Steel drum
Drumstick
Mud turtle

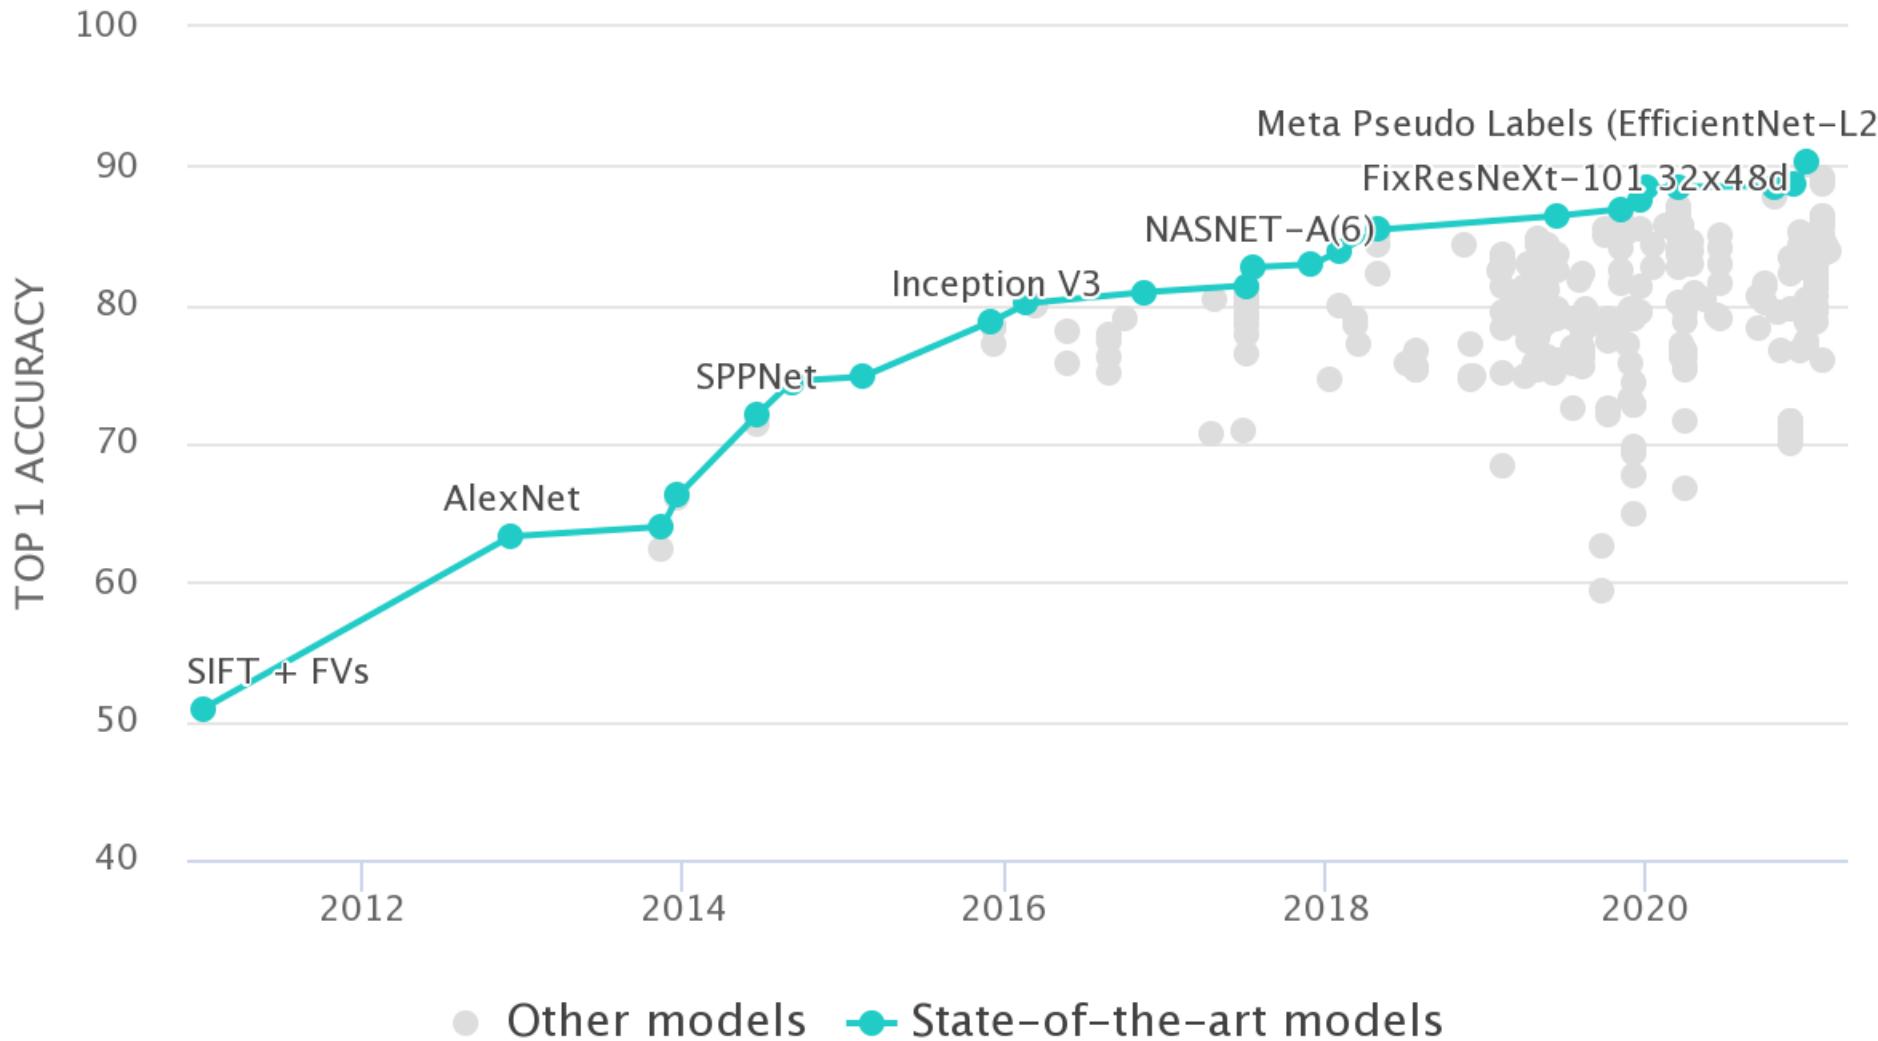


Output:
Scale
Giant panda
Drumstick
Mud turtle



Russakovsky et al. IJCV 2015

IMAGENET Large Scale Visual Recognition Challenge



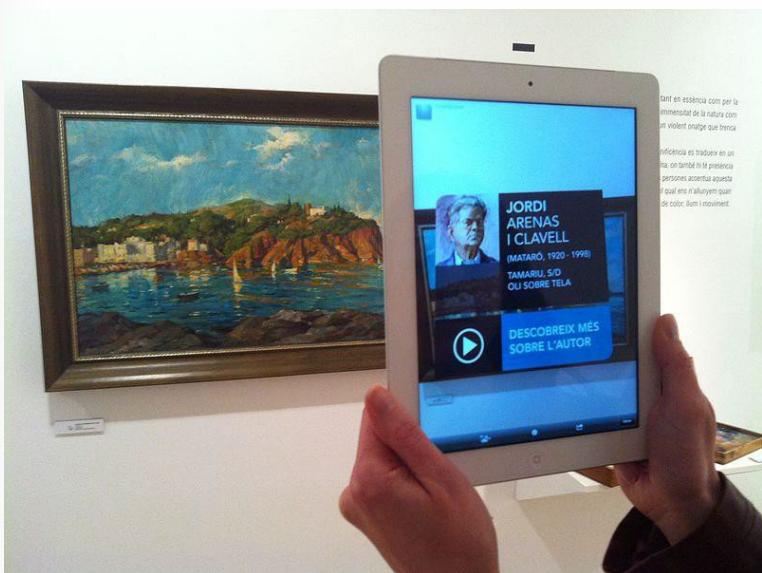
Russakovsky et al. IJCV 2015



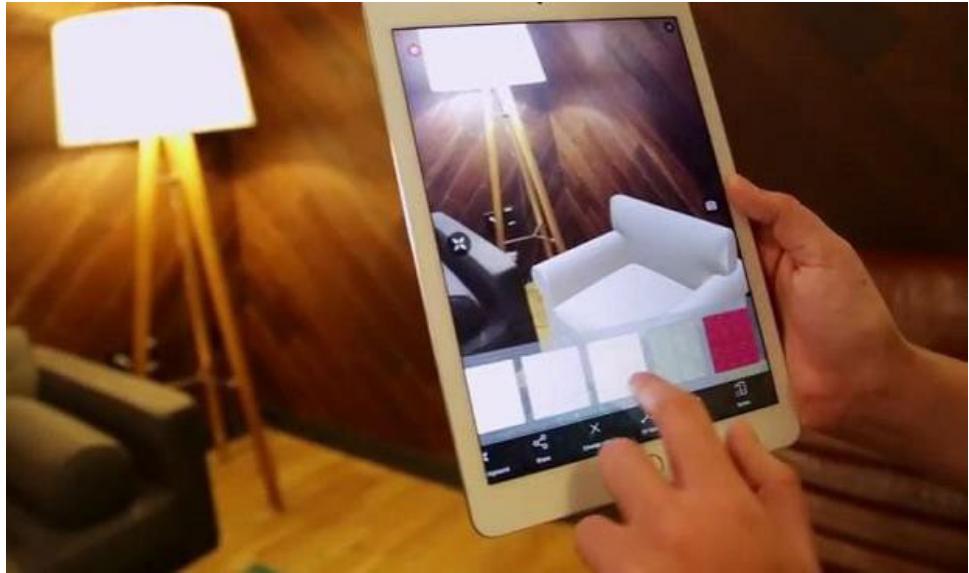
[Image by US Army](#) is licensed under [CC BY 2.0](#)



[Image](#) is [CC0 1.0](#) public domain



[Image by Kippelboy](#) is licensed under [CC BY-SA 3.0](#)

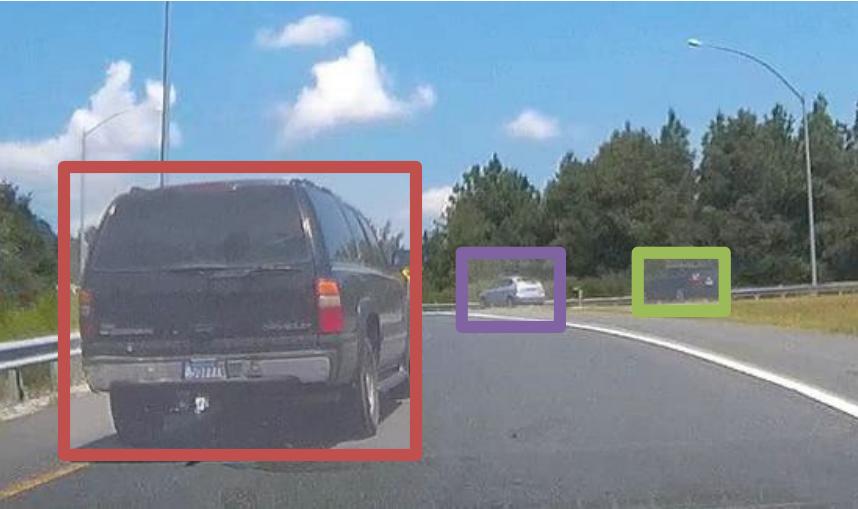


[Image by Christina C.](#) is licensed under [CC BY-SA 4.0](#)

There are many visual recognition problems
that are related to image
classification, such as

object detection, image captioning





[This image](#) is licensed under [CC BY-NC-SA 2.0](#); changes made

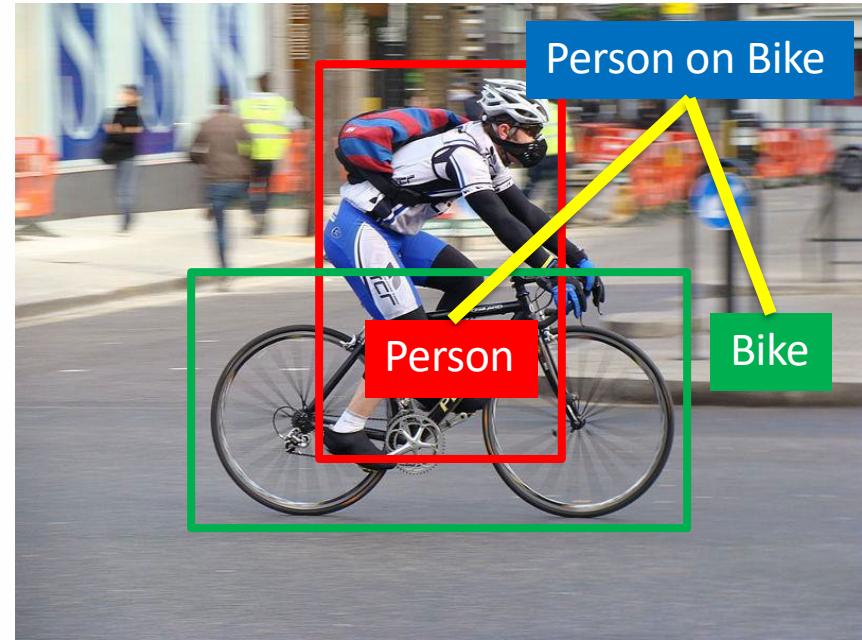


Person

Hammer

[This image](#) is licensed under [CC BY-SA 2.0](#); changes made

- Object detection
- Action classification
- Image captioning
- ...



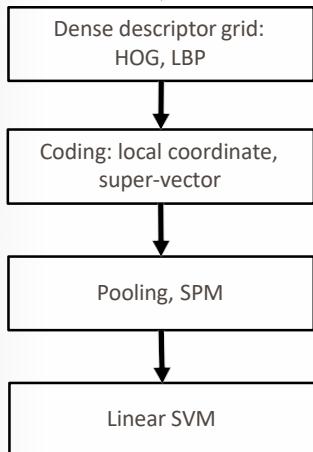
[This image](#) is licensed under [CC BY-SA 3.0](#); changes made

*Convolutional Neural Networks (CNN) have
become an important tool for object recognition*

IMAGENET Large Scale Visual Recognition Challenge

Year 2010

NEC-UIUC

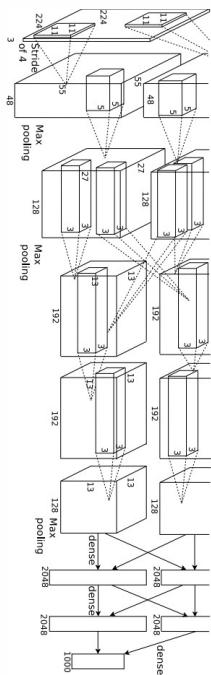


[Lin CVPR 2011]

Lion image by Swissfrog is licensed under CC BY 3.0

Year 2012

SuperVision



[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012
Reproduced with permission.

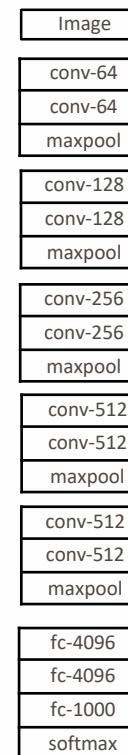
Year 2014

GoogLeNet



[Szegedy arxiv 2014]

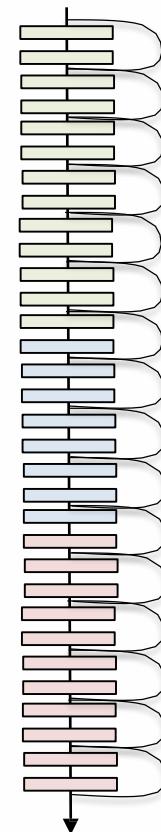
VGG



[Simonyan arxiv 2014]

Year 2015

MSRA

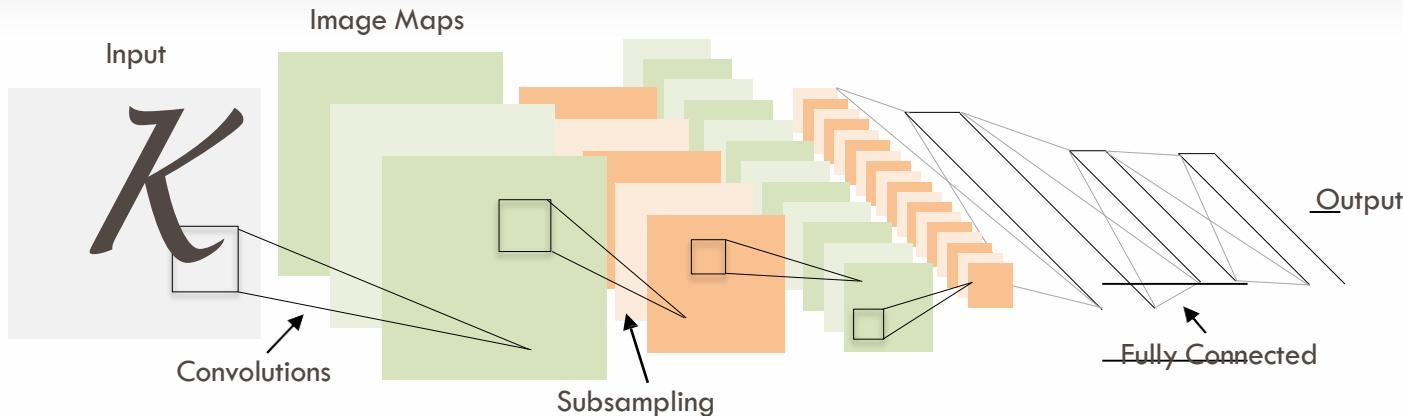


[He ICCV 2015]

Convolutional Neural Networks (CNN)
were not invented overnight

1998

LeCun et al.



of transistors



10^6

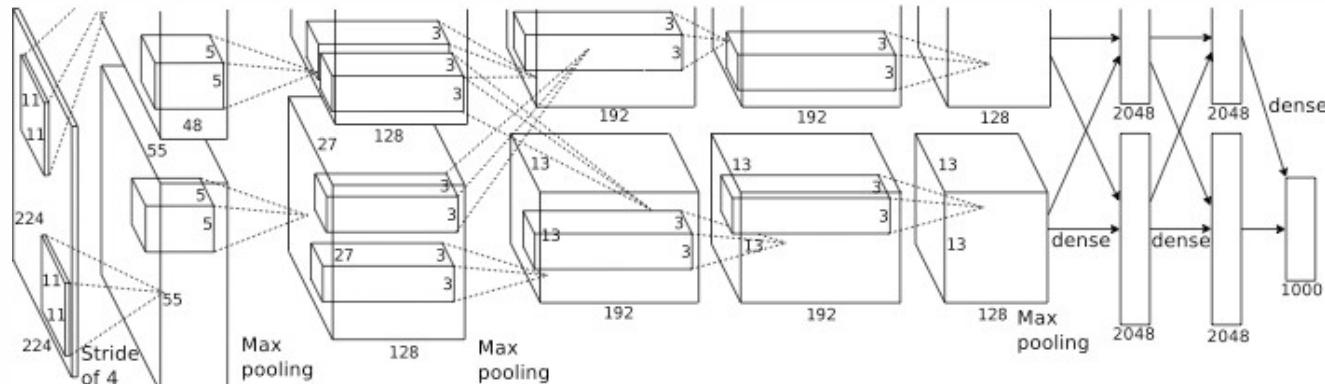
pentium® II

of pixels used in training

10^7 **NIST**

2012

Krizhevsky et al.



of transistors



10^9

GPUs



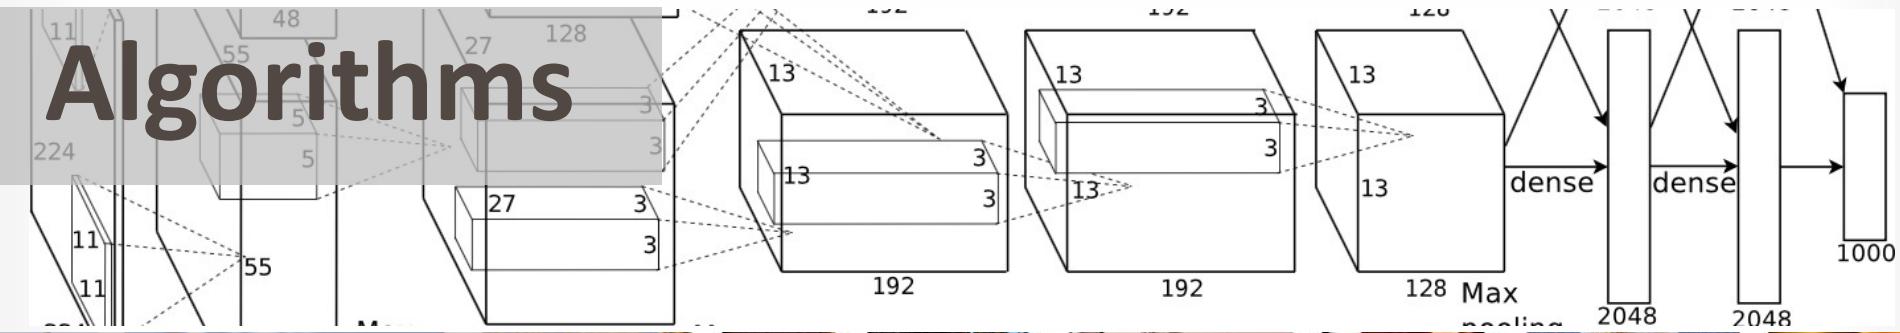
of pixels used in training

10^{14} **IMAGENET**

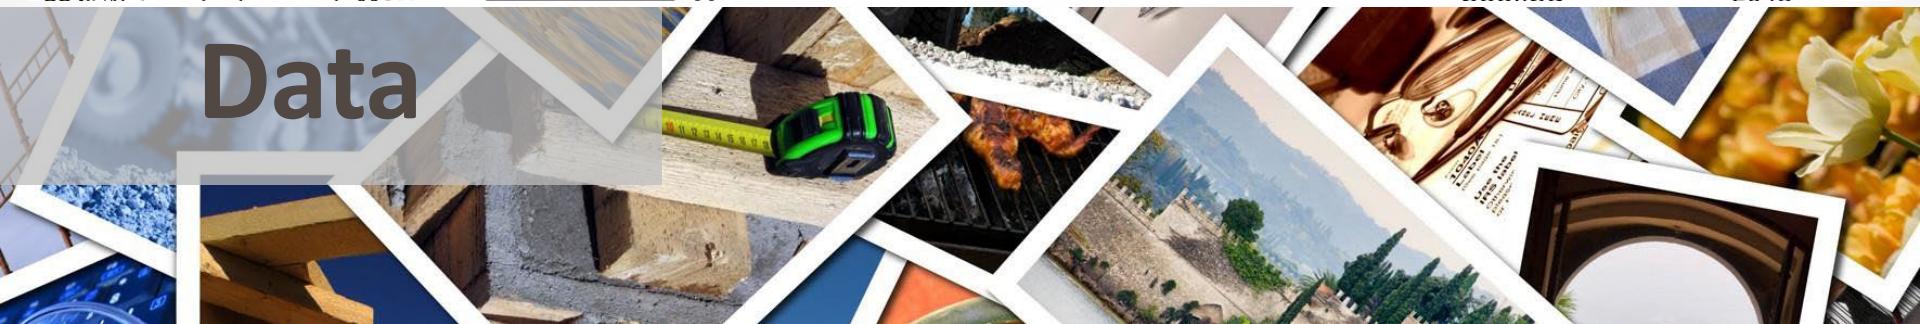
Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

Ingredients for Deep Learning

Algorithms



Data

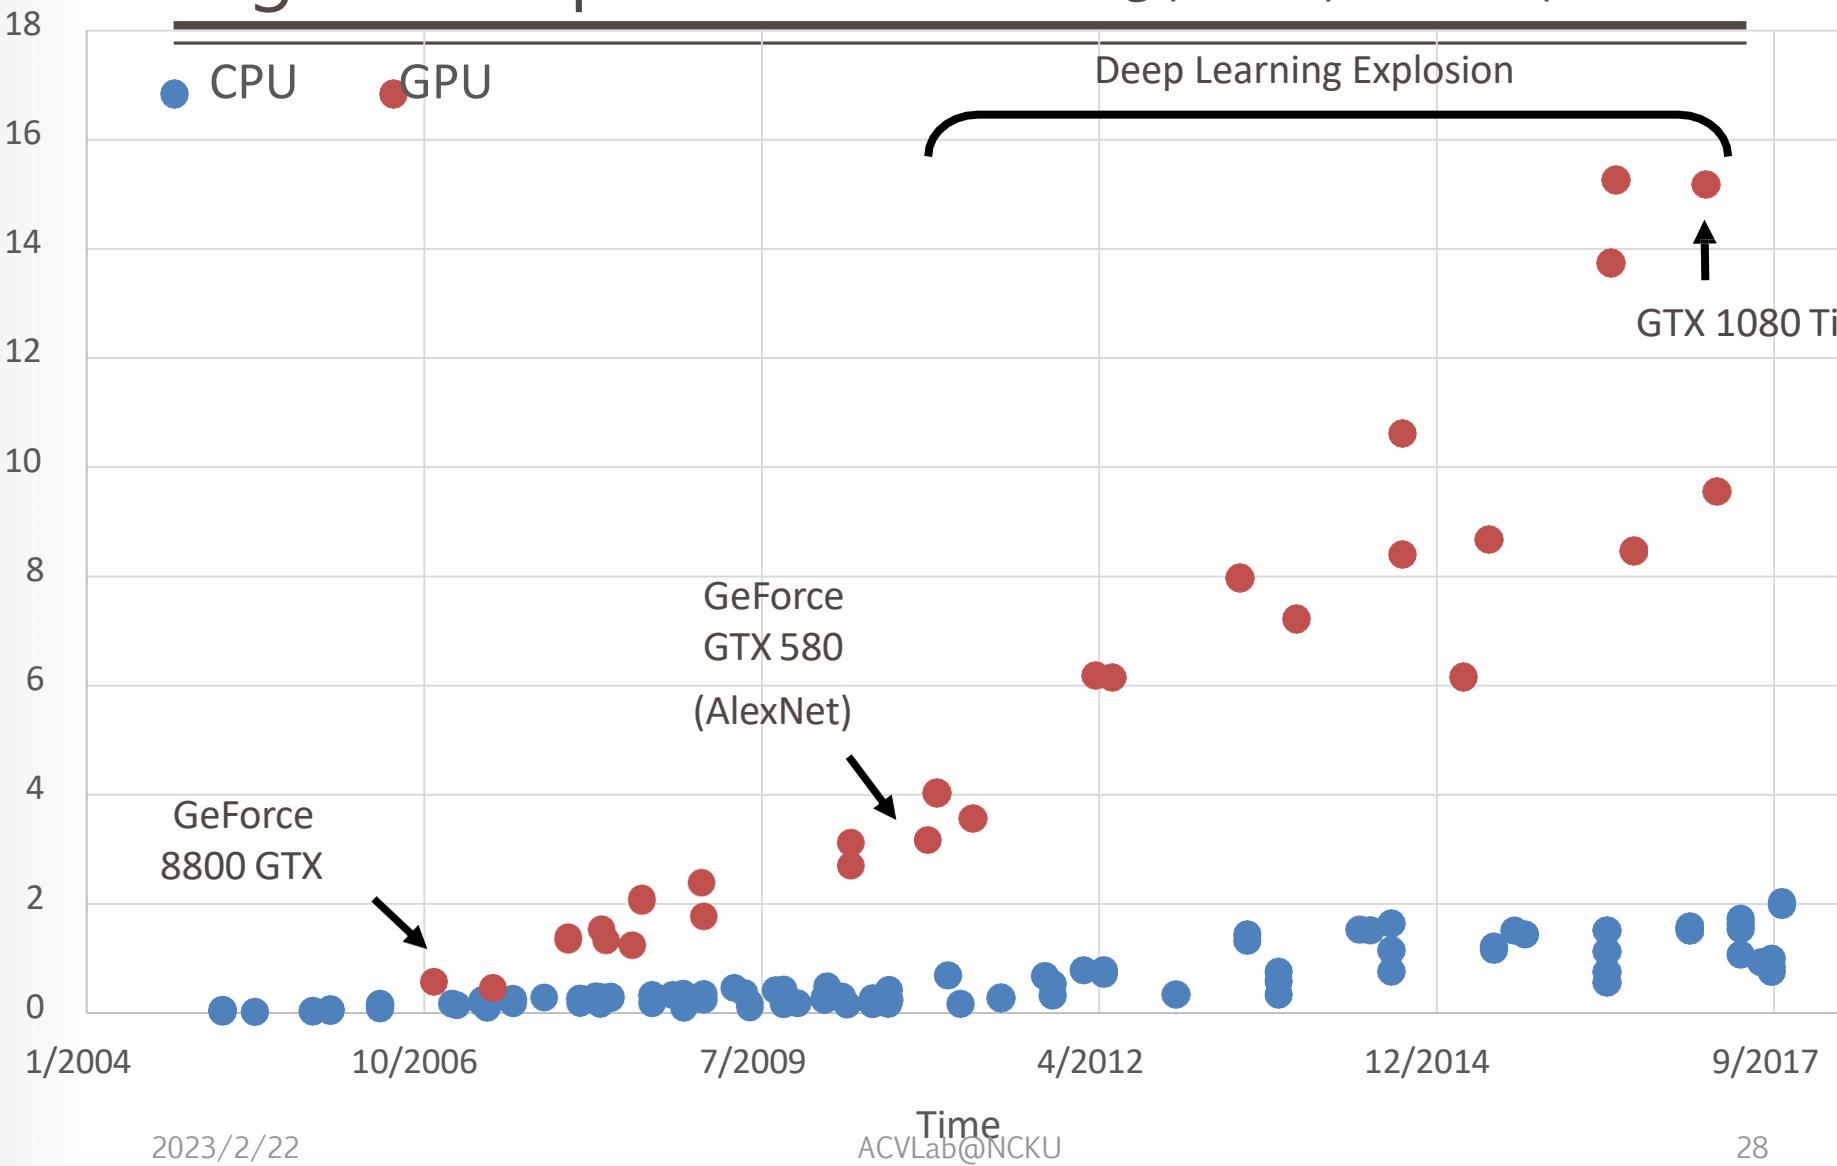


Computation

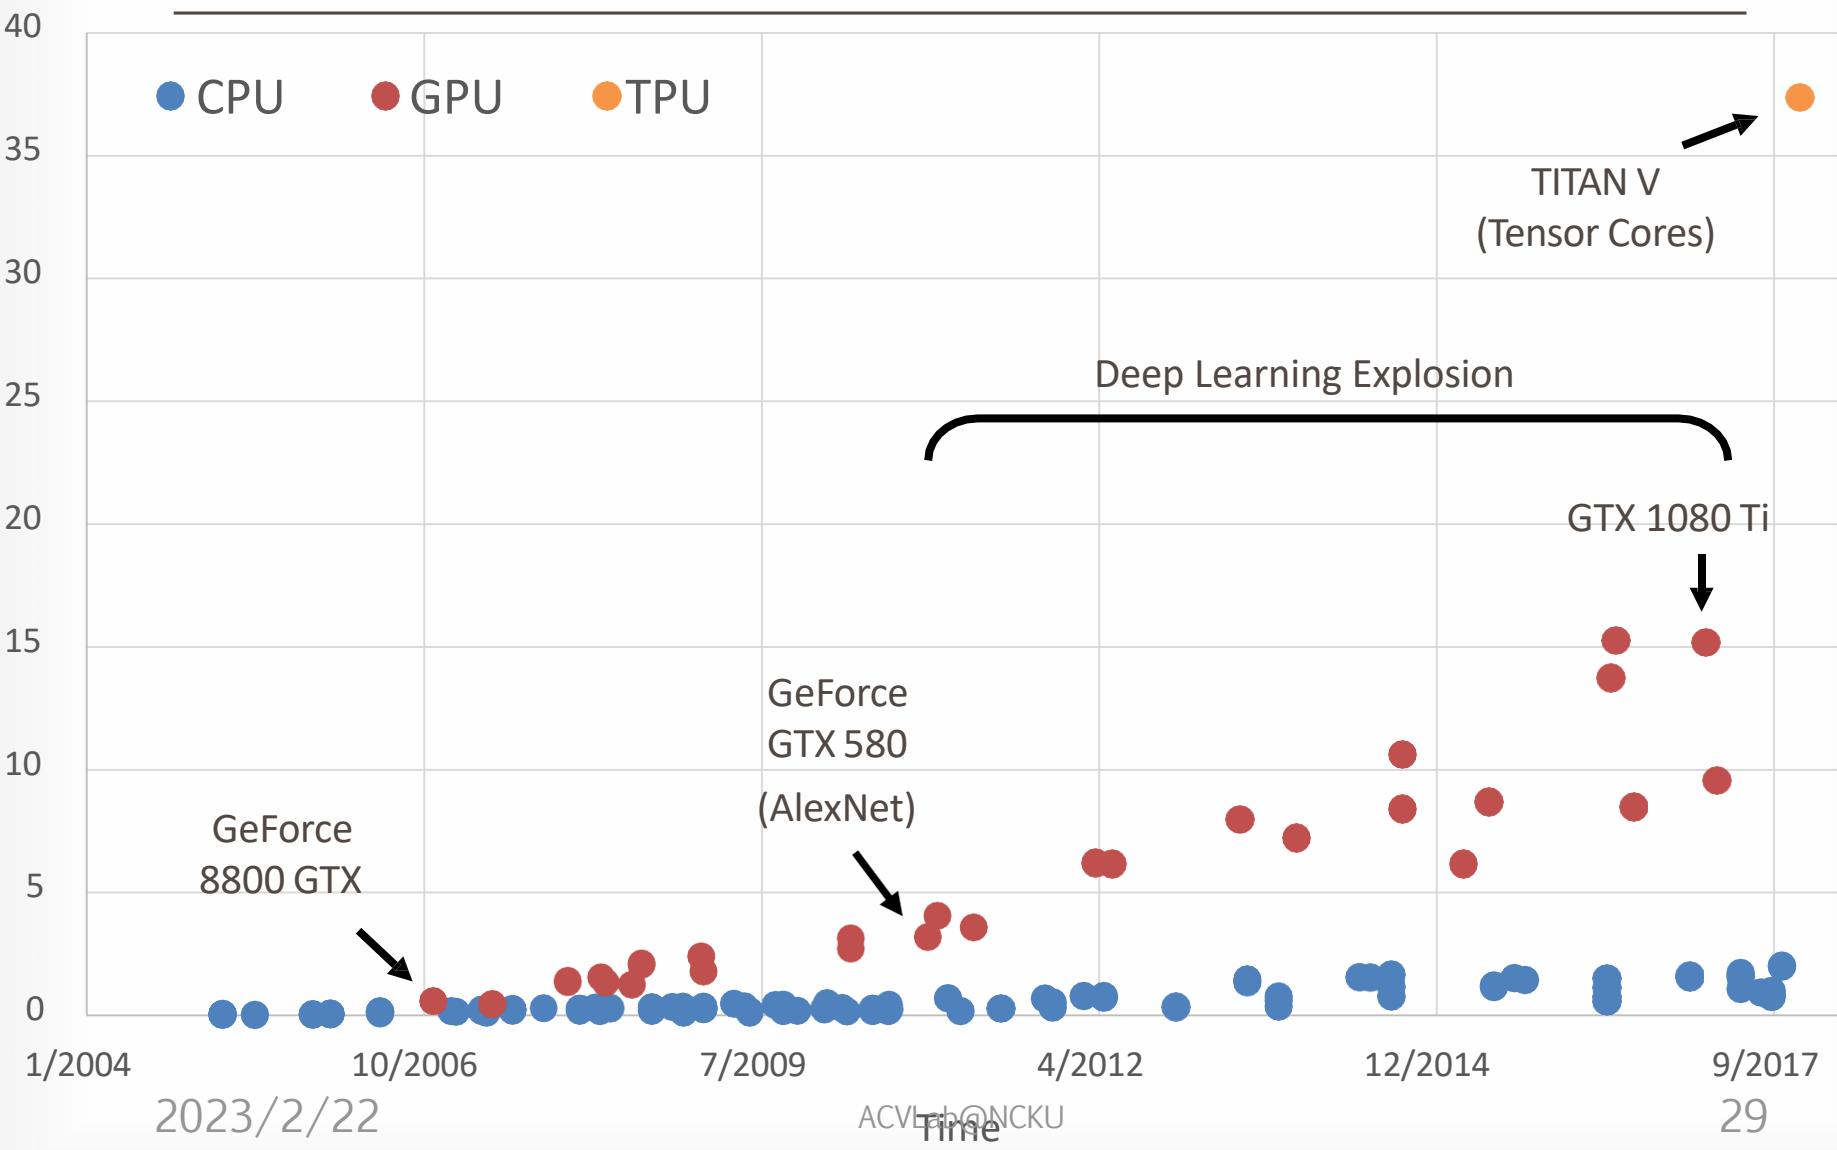


GigaFLOPs per Dollar

Floating-point operations per second



GigaFLOPs per Dollar



The quest for visual intelligence
goes far beyond object recognition...

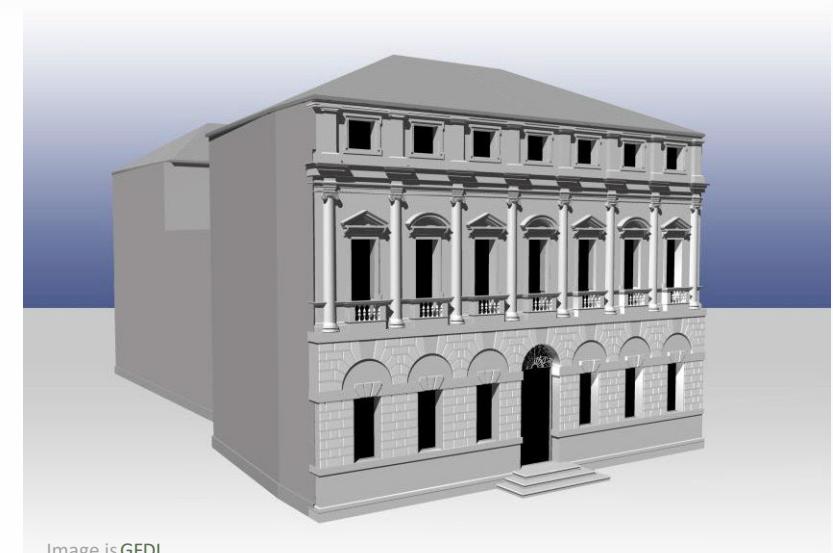
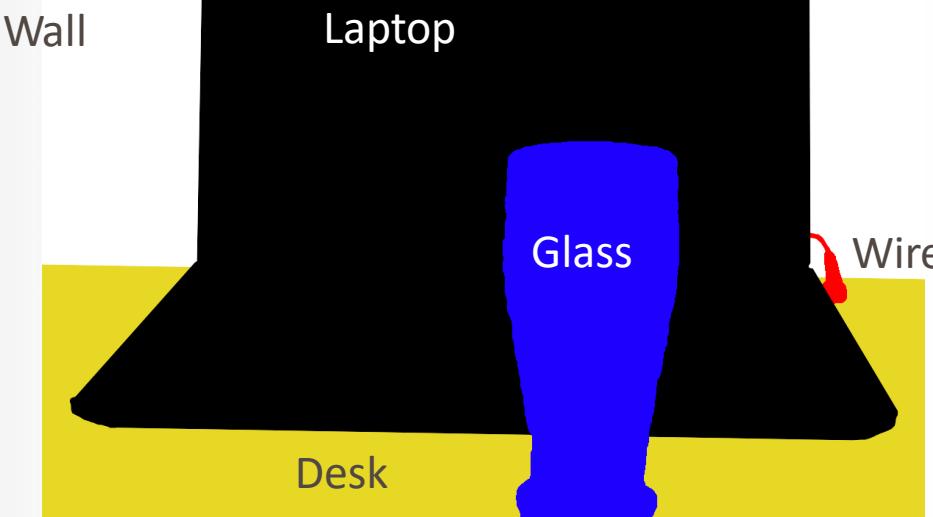
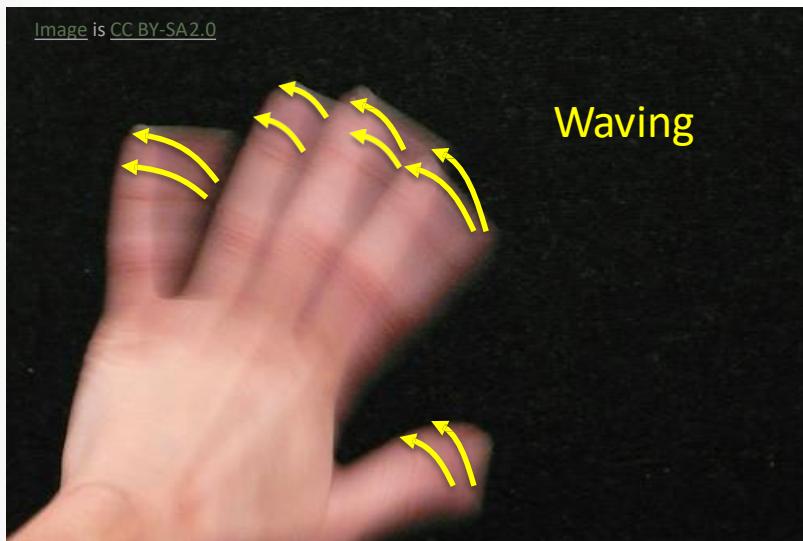
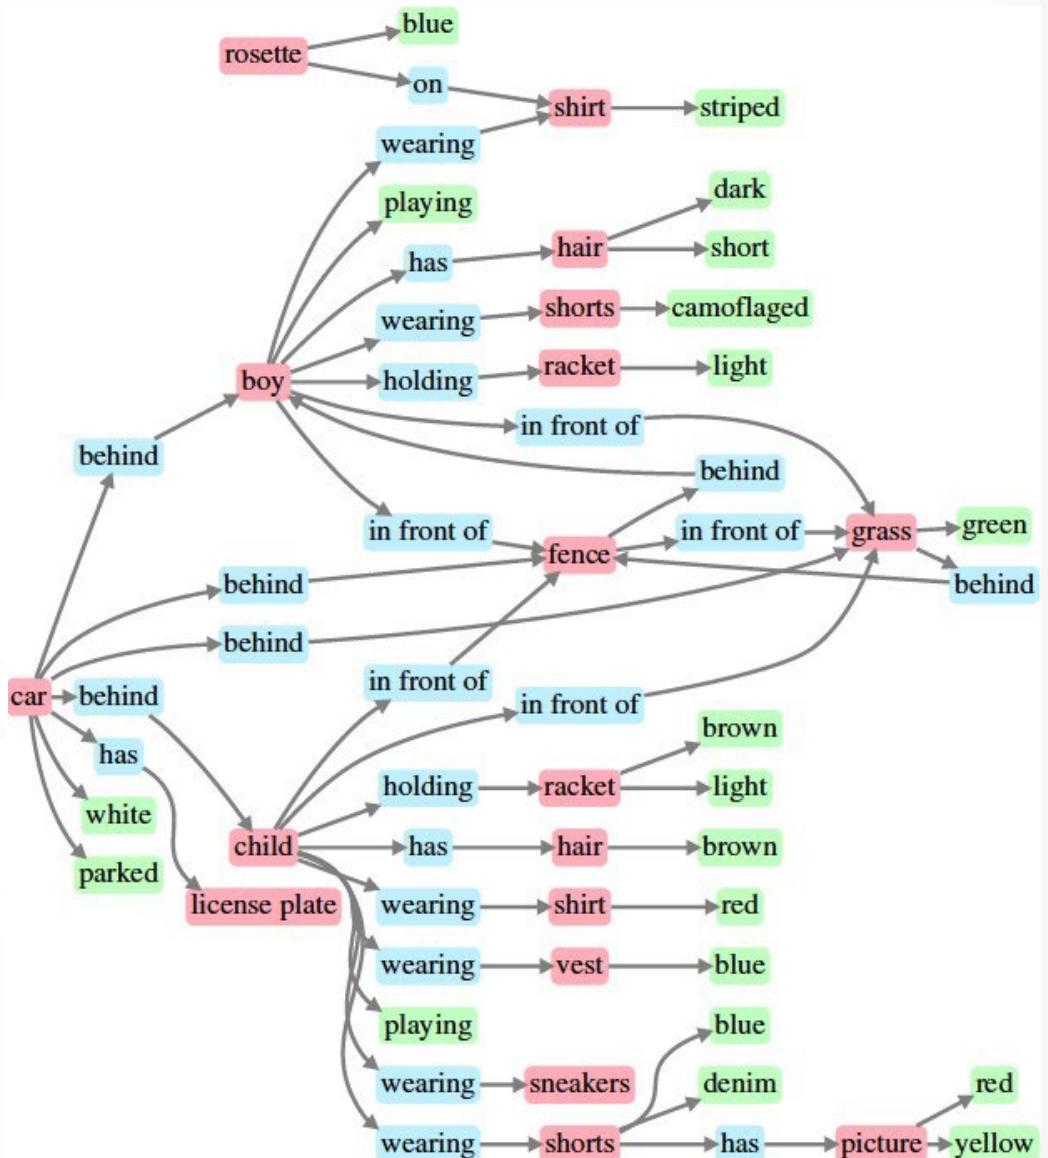
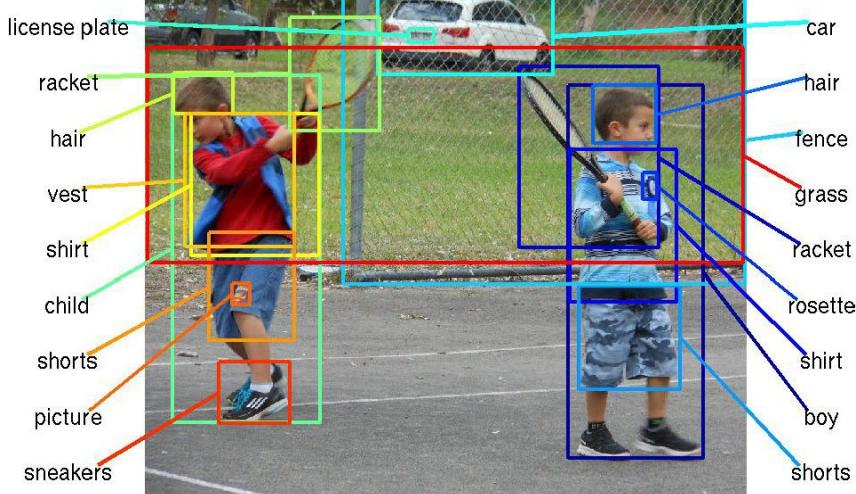


Image is [GFDL](#)





Johnson *et al.*, “Image Retrieval using Scene Graphs”, CVPR 2015

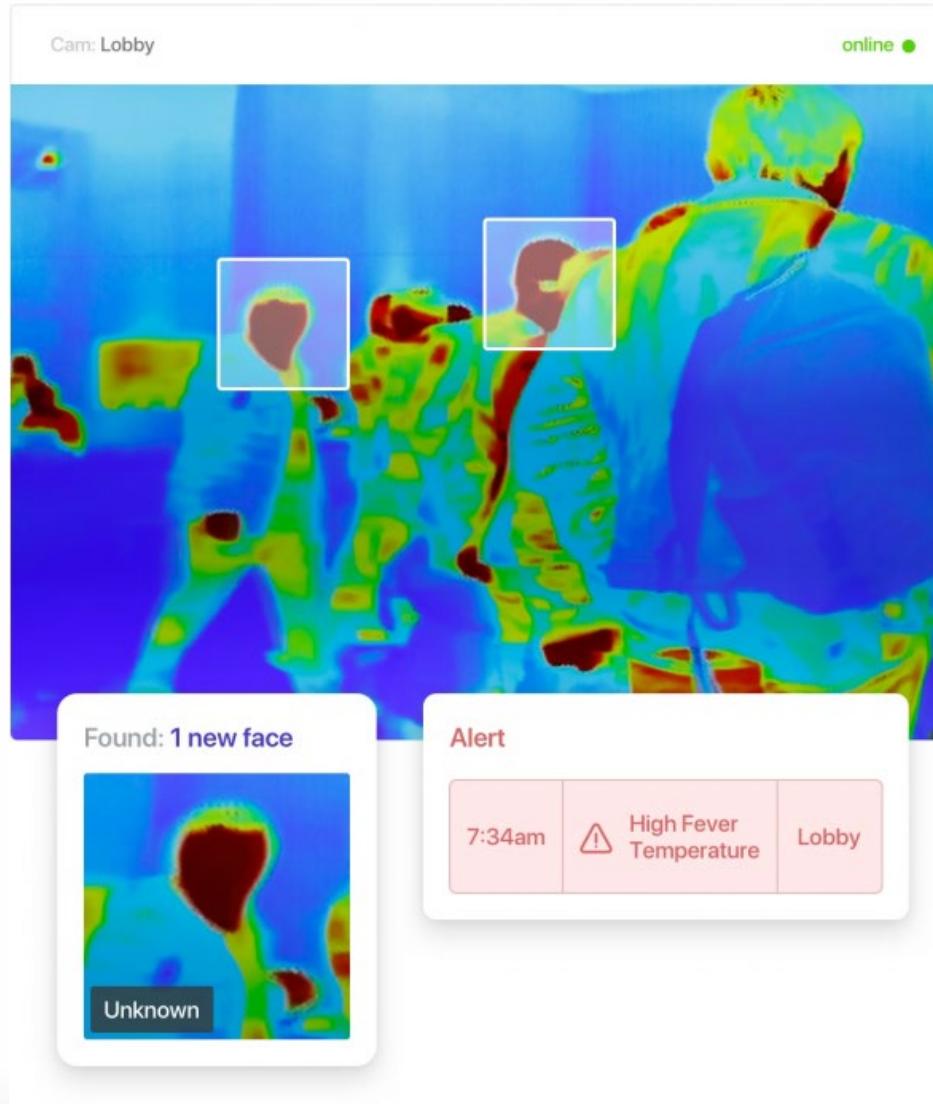
Figures copyright IEEE, 2015. Reproduced for educational purposes



[This image](#) is copyright-free [United States government work](#)

Example credit: [Andrej Karpathy](#)

Face Recognition of NIR Images





Outside border images, clockwise, starting from top left:

[Image by Pop Culture Geek](#) is licensed under [CC BY 2.0](#); changes made
[Image by the US Government](#) is in the public domain
[Image by the US Government](#) is in the public domain
[Image by Glogger](#) is licensed under [CC BY-SA 3.0](#); changes made
[Image by Sylenus](#) is licensed under [CC BY 3.0](#); changes made
[Image by US Government](#) is in the public domain

Inside four images, clockwise, starting from top left:

[Image](#) is [CC0 1.0](#) public domain
[Image by Tucania](#) is licensed under [CC BY-SA 3.0](#); changes made
[Image by Intuitive Surgical, Inc.](#) is licensed under [CC BY-SA 3.0](#); changes made
[Image by Oyundari Zorigtbaatar](#) is licensed under [CC BY-SA 4.0](#)



HOW?



IMAGE CLASSIFICATION PIPELINE

Image Classification: A core task in Computer Vision



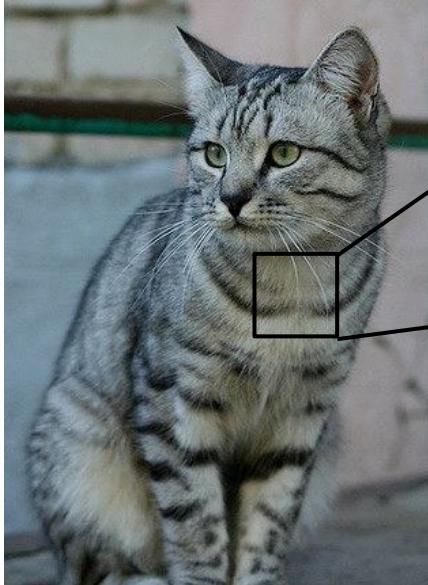
This image by [Nikita](#) is
licensed under [CC-BY 2.0](#)

(assume given set of discrete labels)
{dog, cat, truck, plane, ...}



cat

The Problem: Semantic Gap



This image by [Nikita](#) is
licensed under [CC-BY 2.0](#).

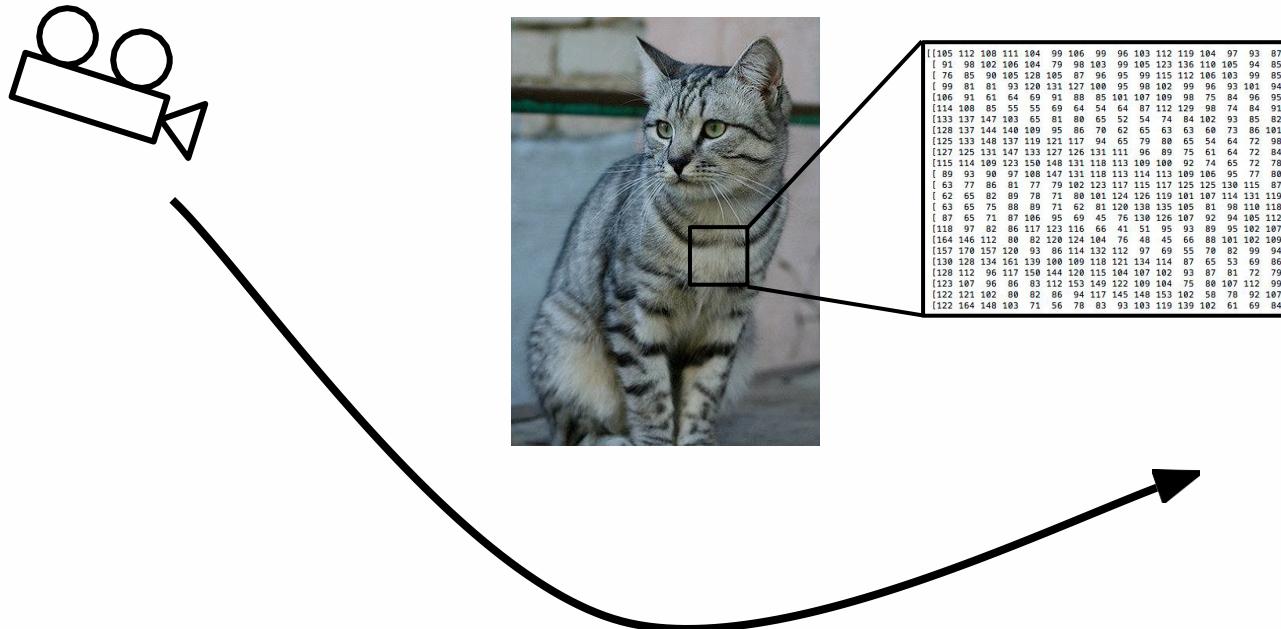
[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
[91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
[76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
[99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
[89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 88]
[63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
[62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
[63 65 75 88 89 71 62 81 120 138 135 105 81 98 118 118]
[87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 97 89 95 102 107]
[164 146 112 88 82 120 124 104 76 48 45 66 88 101 102 109]
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
[123 107 96 86 83 112 153 149 122 189 104 75 80 107 112 99]
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]

What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Challenges: Viewpoint variation



This image by [Nikita](#) is
licensed under [CC-BY 2.0](#)

Challenges: Background Clutter



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Challenges: Illumination



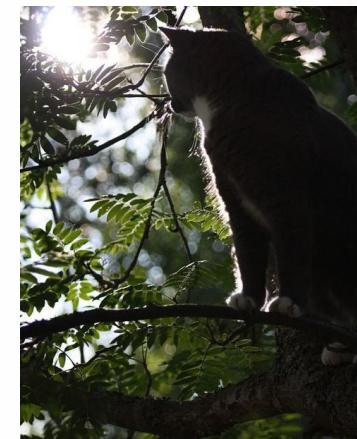
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Challenges: Deformation



[This image](#) by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image](#) by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image](#) by [sare bear](#) is
licensed under [CC-BY 2.0](#)



[This image](#) by [Tom Thai](#) is
licensed under [CC-BY 2.0](#)

Challenges: Occlusion



[This image](#) is CC0 1.0 public domain



[This image](#) is CC0 1.0 public domain



[This image](#) by [jonsson](#) is licensed under CC-BY 2.0

Challenges: Intraclass variation



[This image](#) is CC0 1.0 public domain

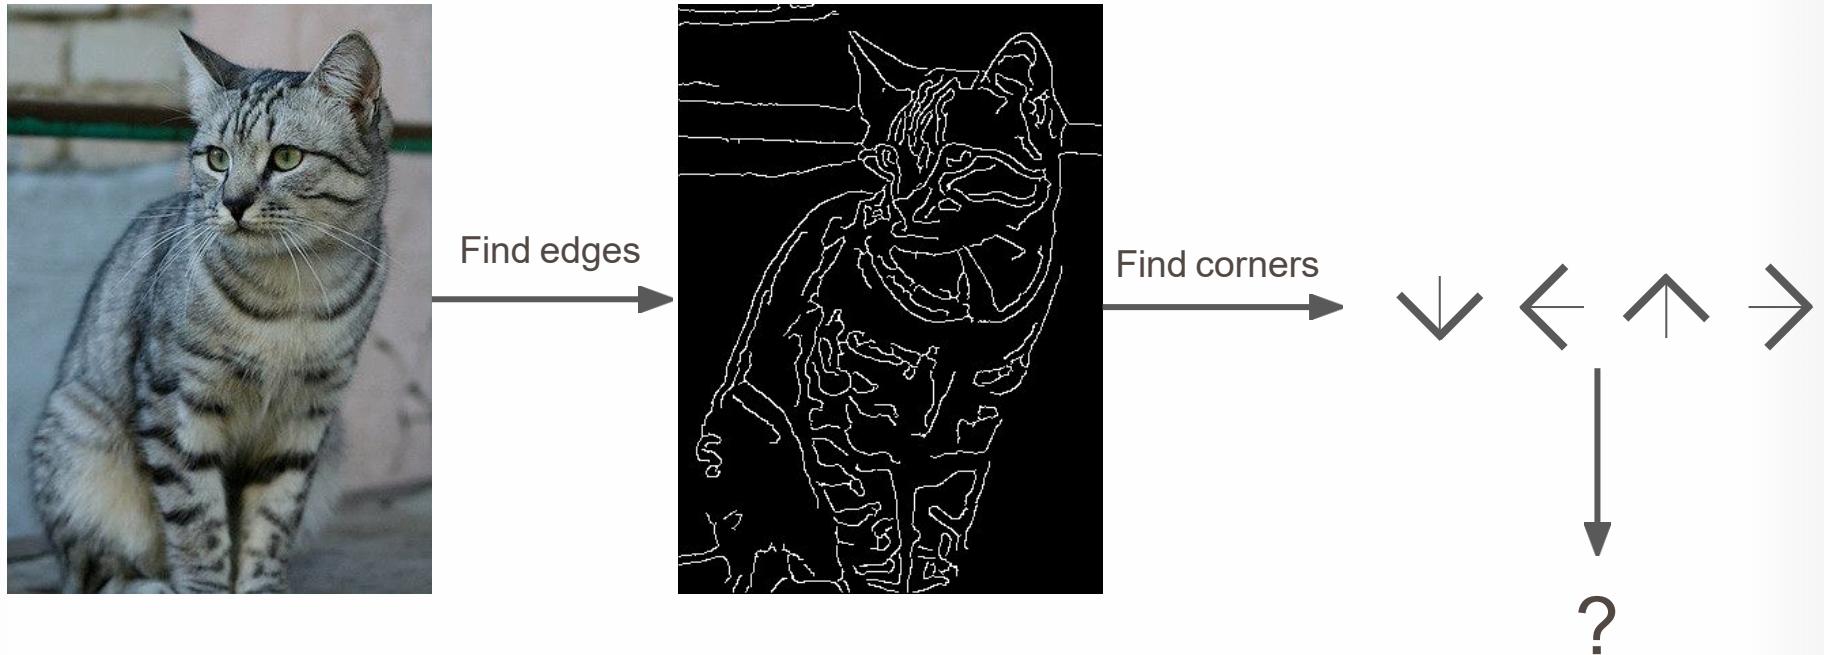
An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for
recognizing a cat, or other classes.

Attempts have been made



John Canny, "A Computational Approach to Edge Detection", IEEE TPAMI 1986

Machine Learning: Data-Driven Approach

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

Example training set

```
def train(images, labels):  
    # Machine learning!  
    return model
```

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

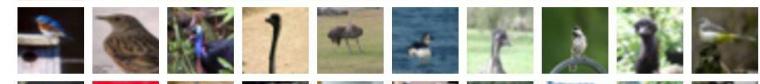
airplane



automobile



bird



cat



deer



First classifier: Nearest Neighbor

```
def train(images, labels):  
    # Machine learning!  
    return model
```



Memorize all
data and labels

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```



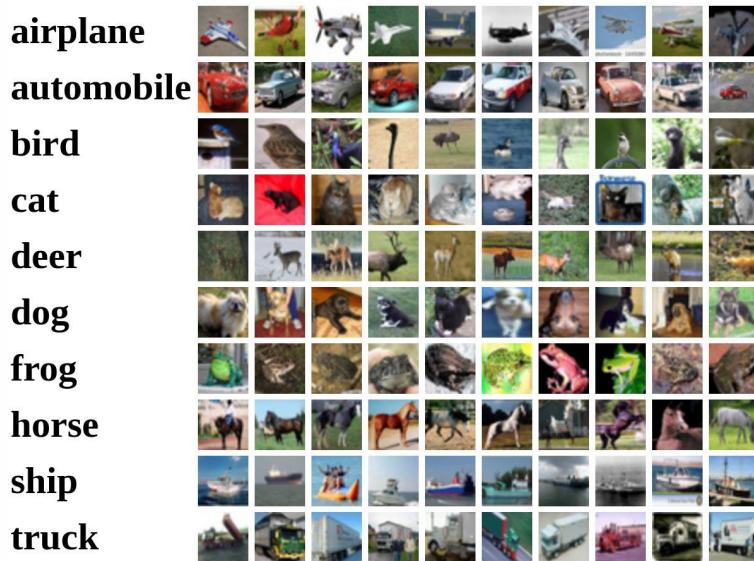
Predict the label
of the most similar
training image

Example Dataset: CIFAR10

10 classes

50,000 training images

10,000 testing images



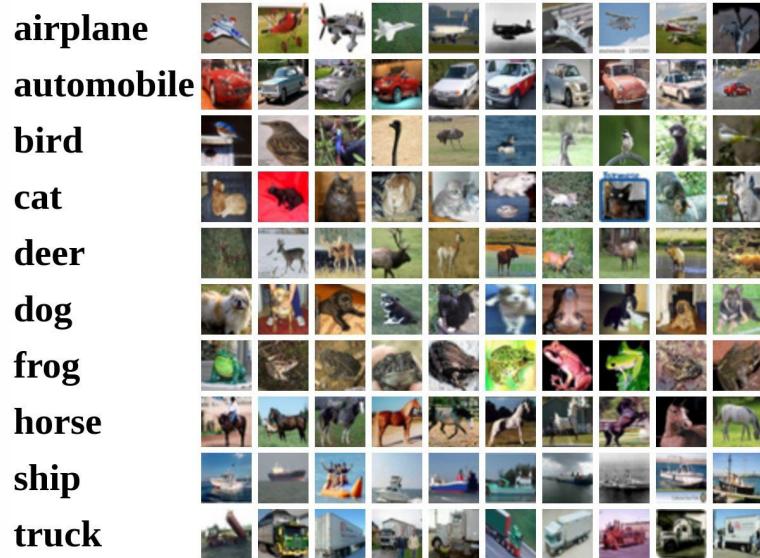
Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", Technical Report, 2009.

Example Dataset: CIFAR10

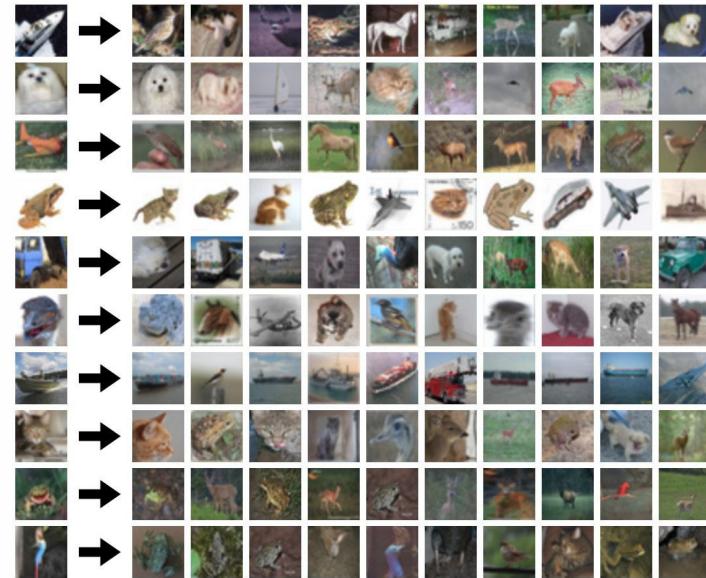
10 classes

50,000 training images

10,000 testing images



Test images and nearest neighbors



Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", Technical Report, 2009.

Distance Metric to compare images

L1 distance:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

- = add → 456

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

Memorize training data

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

For each test image:
 Find closest train image
 Predict label of nearest image

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

Q: With N examples, how fast are training and prediction?

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

Q: With N examples, how fast are training and prediction?

A: Train $O(1)$, predict $O(N)$

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

Q: With N examples, how fast are training and prediction?

A: Train $O(1)$, predict $O(N)$

This is bad: we want classifiers that are **fast** at prediction; **slow** for training is ok

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor classifier

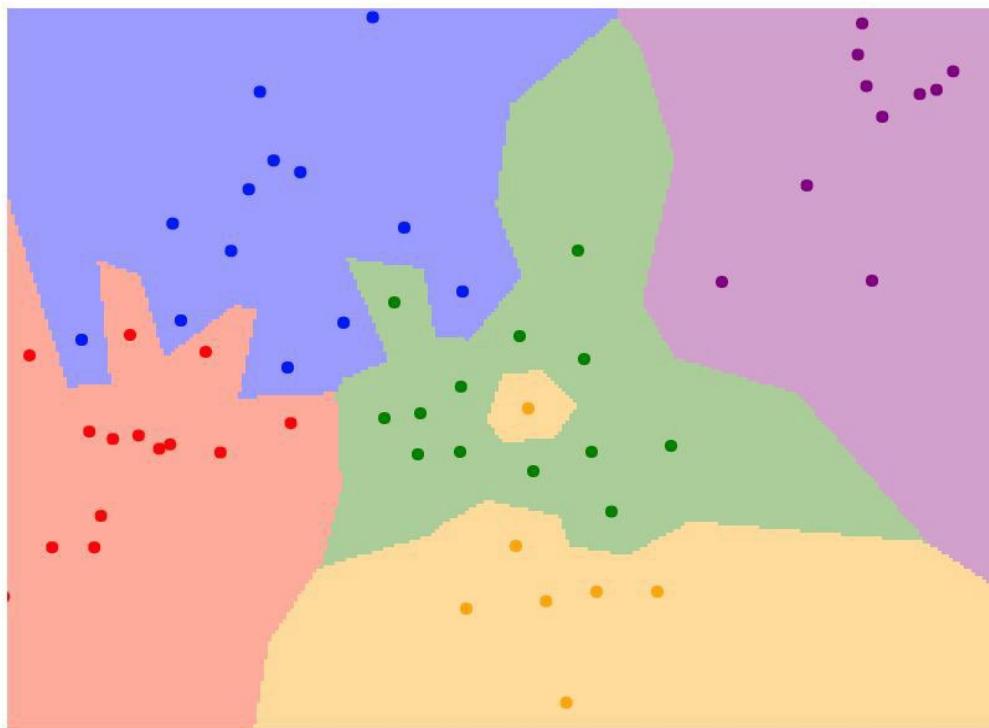
Many methods exist for fast / approximate nearest neighbor

A good implementation:

<https://github.com/facebookresearch/faiss>

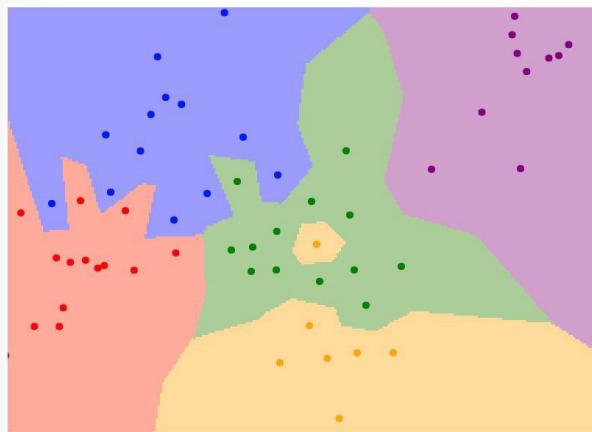
Johnson et al, “Billion-scale similarity search with GPUs”, arXiv 2017

What does this look like?

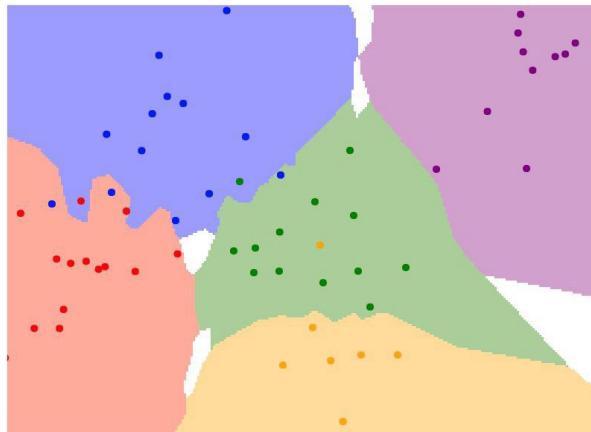


K-Nearest Neighbors

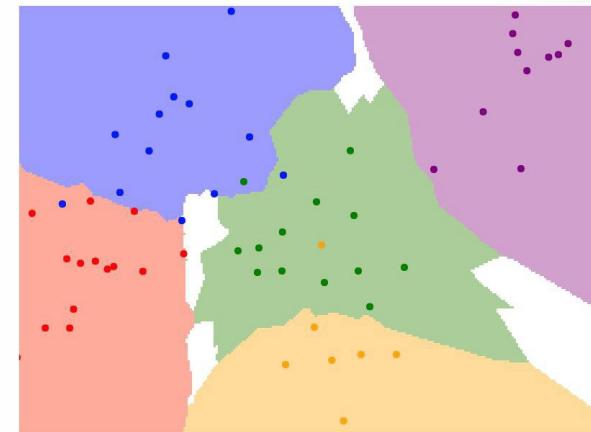
Instead of copying label from nearest neighbor,
take **majority vote** from K closest points



$K = 1$

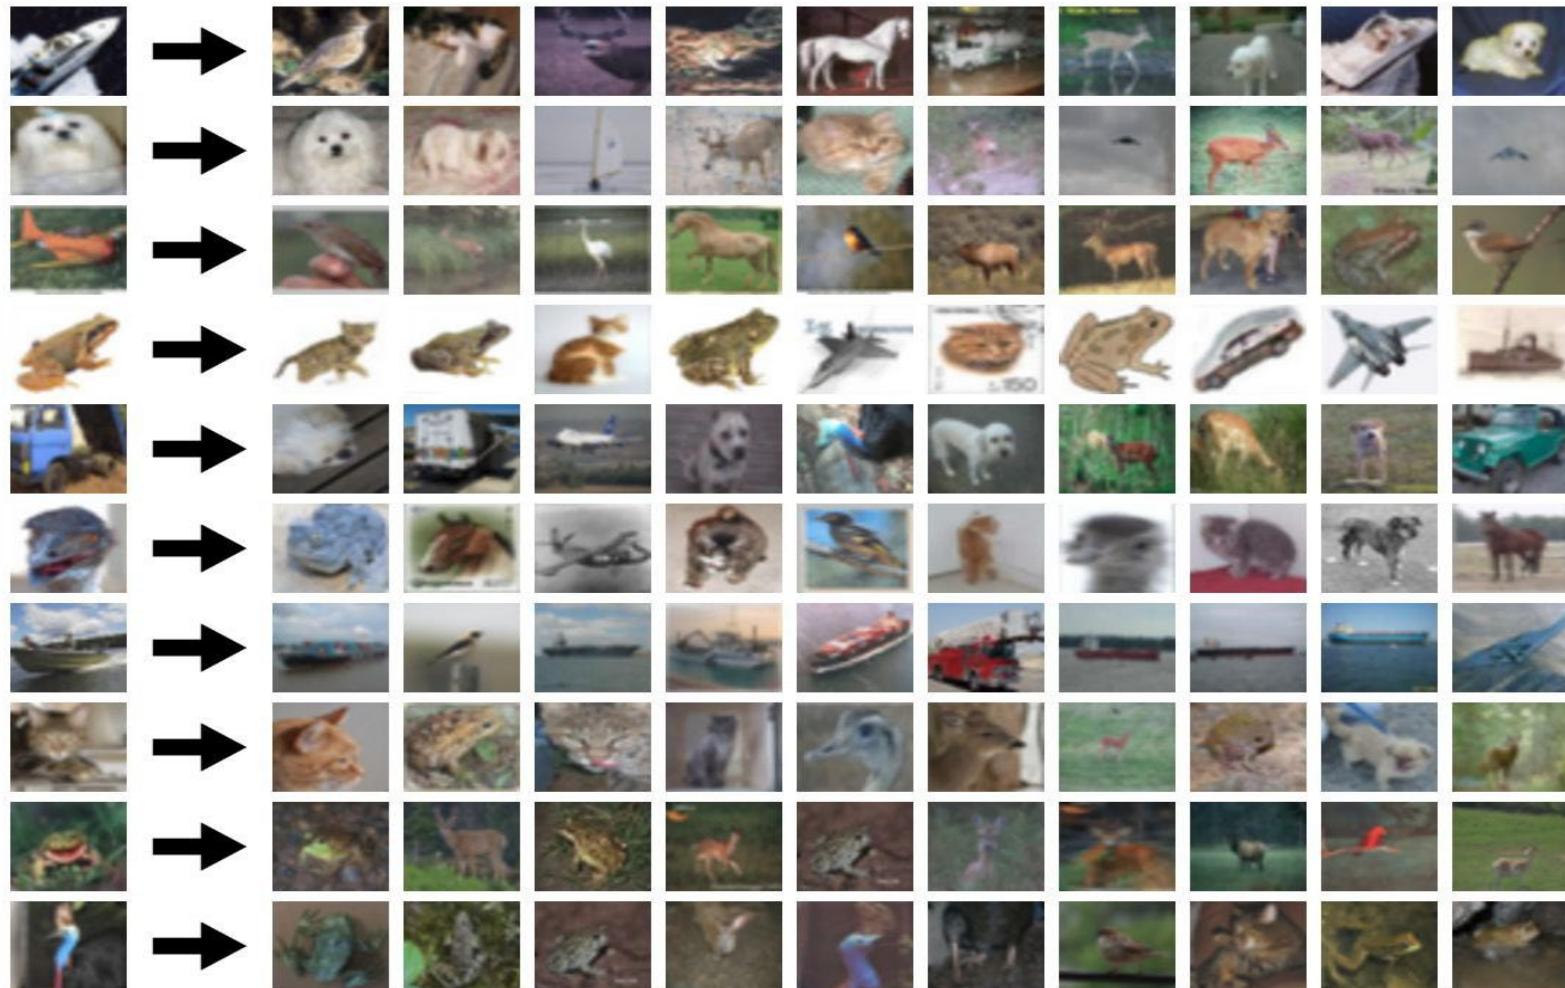


$K = 3$

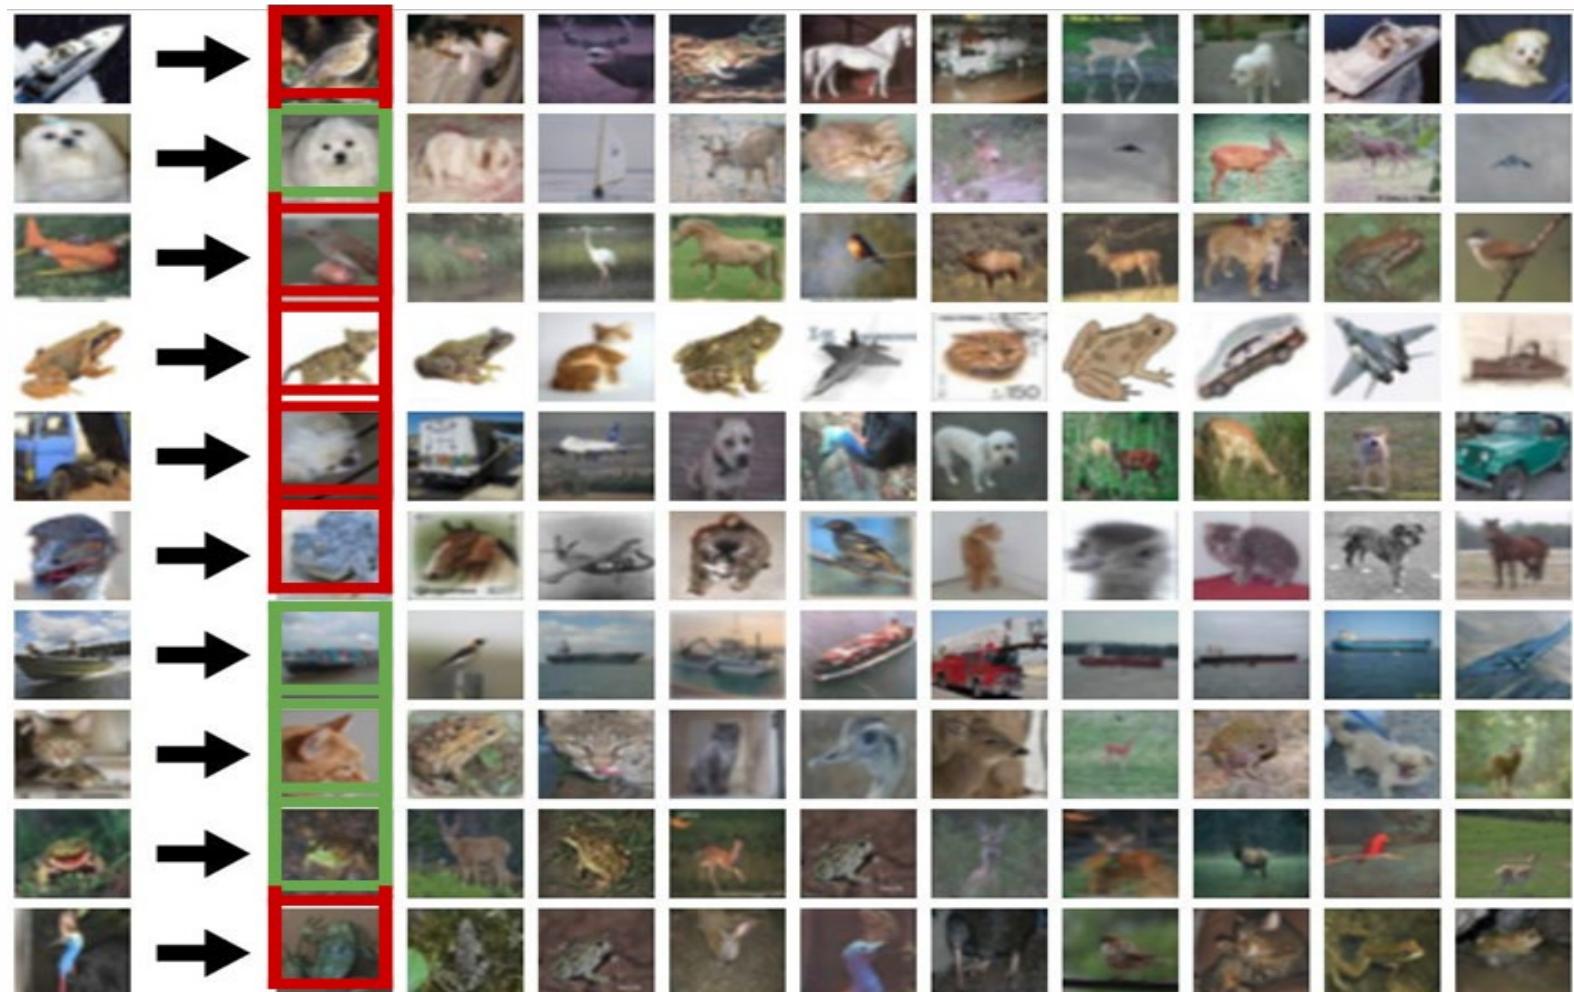


$K = 5$

What does this look like?



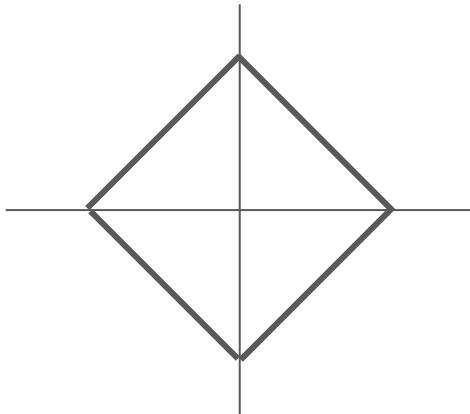
What does this look like?



K-Nearest Neighbors: Distance Metric

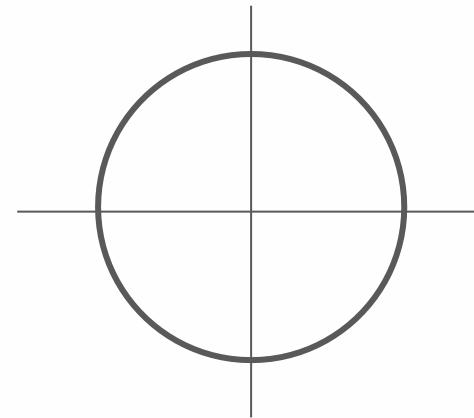
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

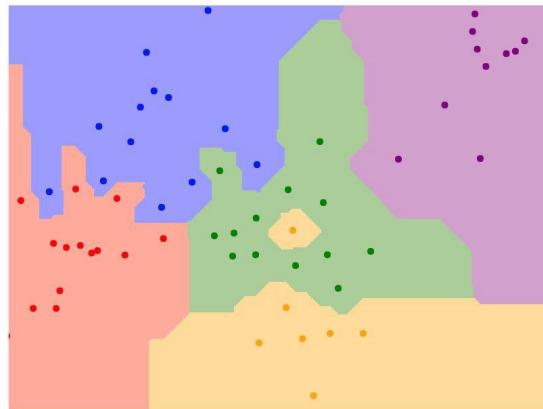
$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



K-Nearest Neighbors: Distance Metric

L1 (Manhattan) distance

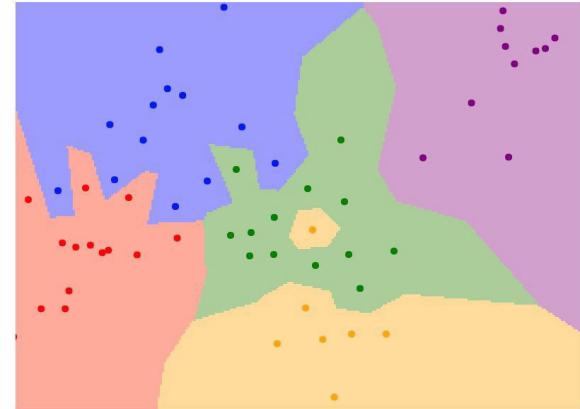
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



$K = 1$

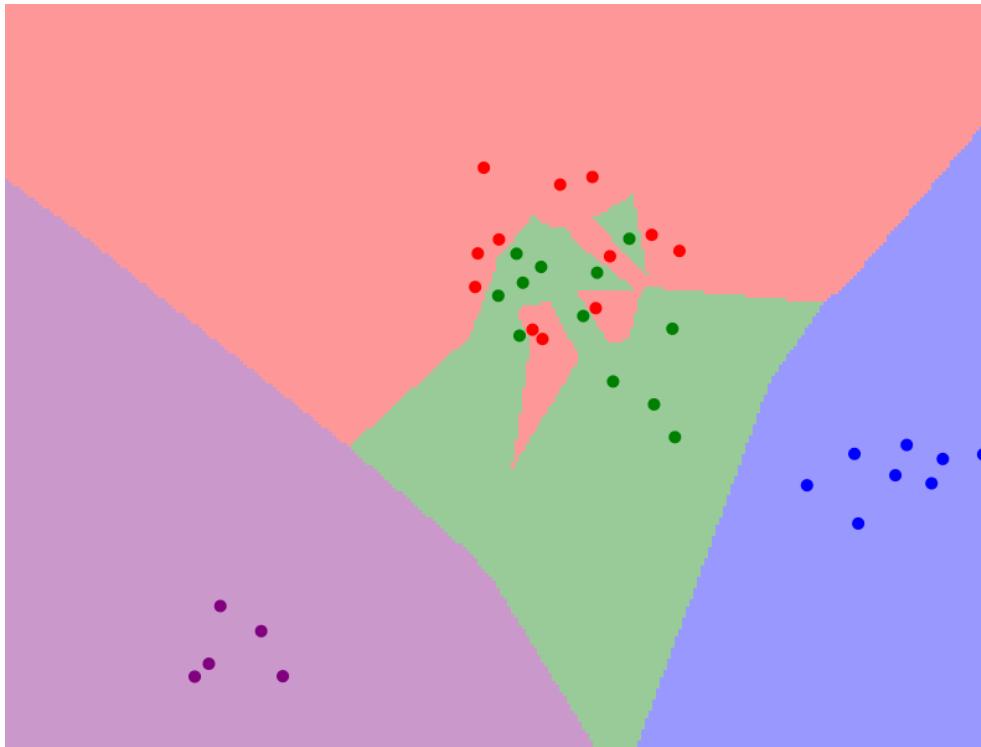
L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



$K = 1$

K-Nearest Neighbors: Demo Time



<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Hyperparameters

- What is the best value of k to use? What is the best distance to use?
- These are hyperparameters: choices about the algorithm that we set rather than learn

Hyperparameters

- What is the best value of k to use? What is the best distance to use?
- These are hyperparameters: choices about the algorithm that we set rather than learn
- Very problem-dependent.
- Must try them all out and see what works best.

Setting Hyperparameters: Evaluation

Idea #1: Choose hyperparameters
that work best on the data

Your Dataset

Setting Hyperparameters: Evaluation

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

Your Dataset

Setting Hyperparameters: Evaluation

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

train

test

Setting Hyperparameters: Evaluation

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

train

BAD: No idea how algorithm will perform on new data

test

Setting Hyperparameters: Evaluation

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

train

validation

test

Setting Hyperparameters: Evaluation

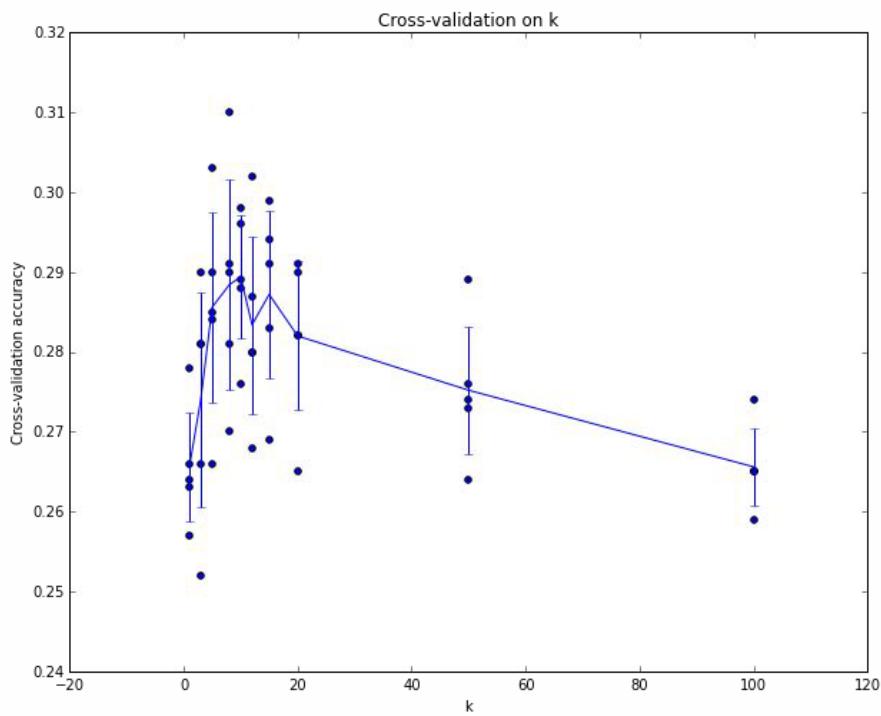
Your Dataset

Idea #4: Cross-Validation: Split data into **folds**,
try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

Setting Hyperparameters: Evaluation



Example of
5-fold cross-validation
for the value of k .

Each point: single
outcome.

The line goes
through the mean, bars
indicated standard
deviation

(Seems that $k \approx 7$ works best
for this data)

k-Nearest Neighbor on images never used.

- Very slow at test time
- Distance metrics on pixels are not informative

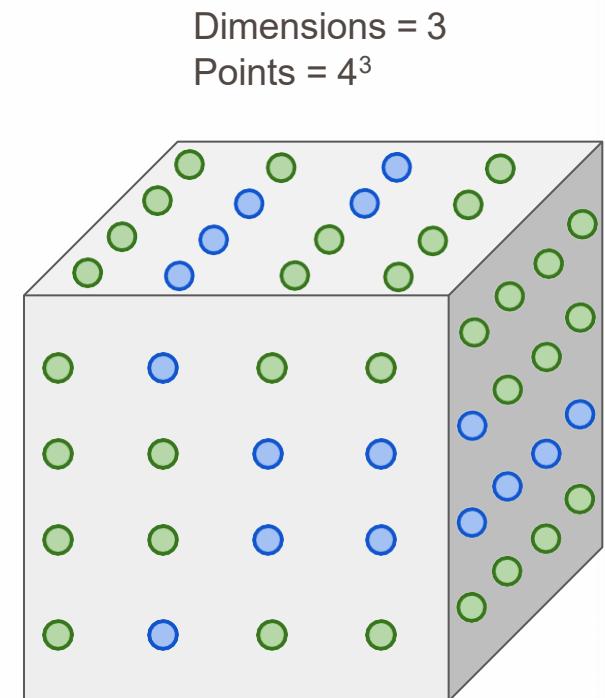
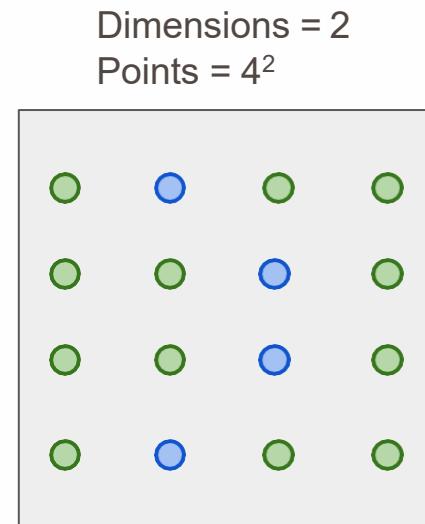
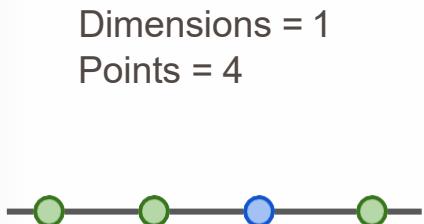


(all 3 images have same L2 distance to the one on the left)

Original image is
CC0 public domain

k-Nearest Neighbor on images never used.

- Curse of dimensionality



k-Nearest Neighbors: Summary

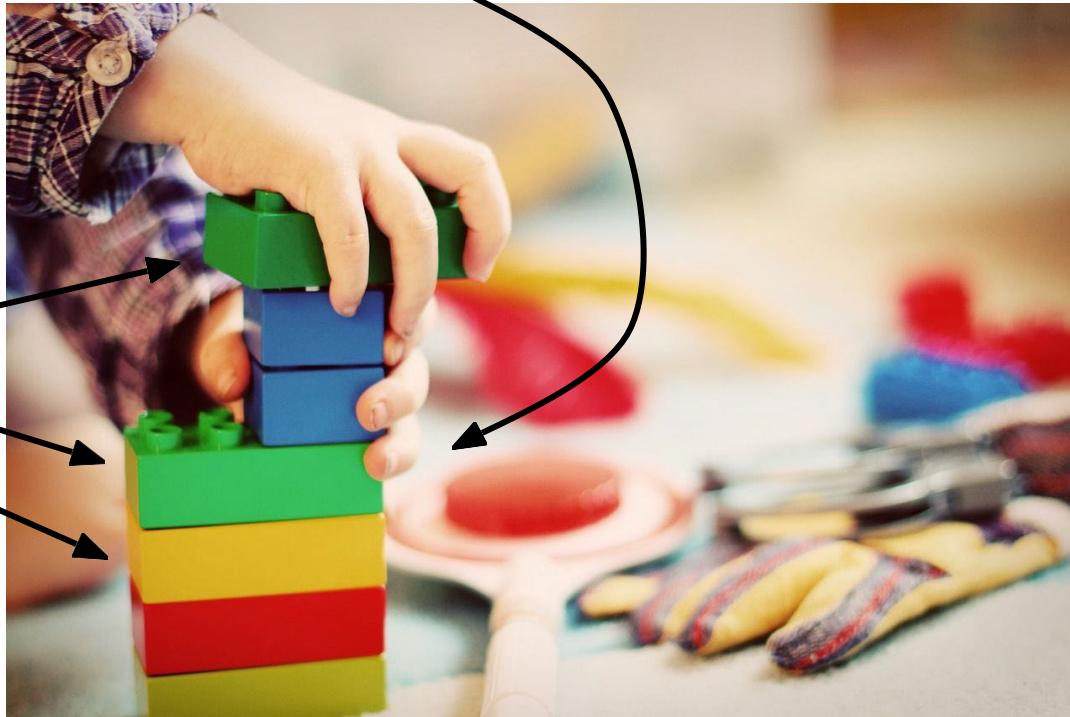
- In Image classification we start with a training set of images and labels, and must predict labels on the test set
- The K-Nearest Neighbors classifier predicts labels based on nearest training examples
- Distance metric and K are **hyperparameters**
- Choose hyperparameters using the validation set; only run on the test set once at the very end!



LINEAR CLASSIFICATION

Neural Network

Linear
classifiers



[This image](#) is CC0 1.0 public domain

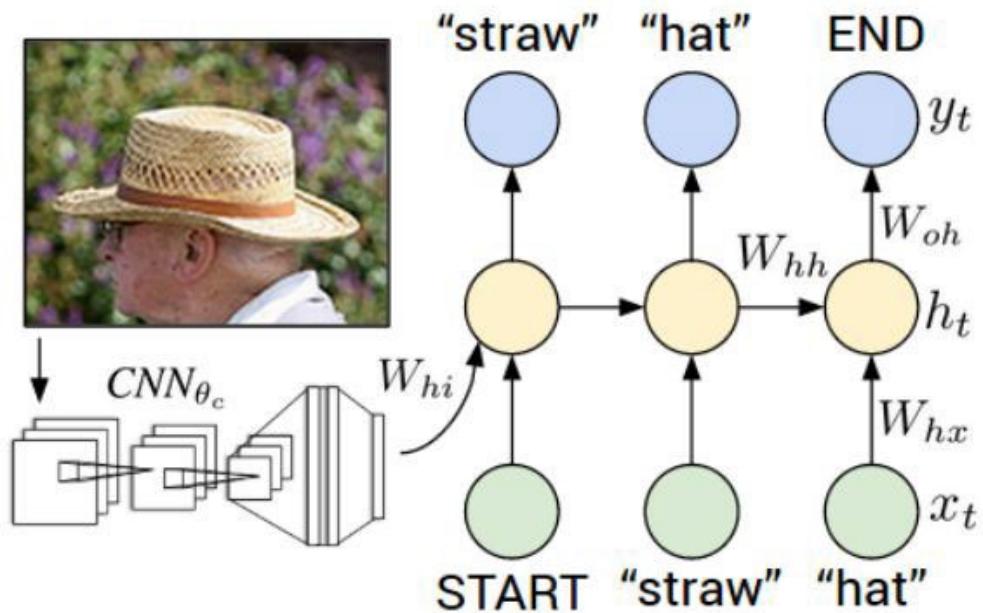
Two young girls are playing with lego toy.



Boy is doing backflip on wakeboard

Man in black shirt is playing guitar.

Construction worker in orange safety vest is working on road.



Karpathy and Fei-Fei, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015
Figures copyright IEEE, 2015. Reproduced for educational purposes.

Recall CIFAR10

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



50,000 training images
each image is **32x32x3**

10,000 test images.

Parametric Approach

Image



Array of **32x32x3** numbers
(3072 numbers total)

$$\xrightarrow{\hspace{1.5cm}} f(\mathbf{x}, \mathbf{W}) \xrightarrow{\hspace{1.5cm}}$$

\mathbf{W}

parameters
or weights

10 numbers giving
class scores

Parametric Approach: Linear Classifier



Array of **32x32x3** numbers
(3072 numbers total)

$$f(\mathbf{x}, \mathbf{W}) = \mathbf{W}\mathbf{x}$$

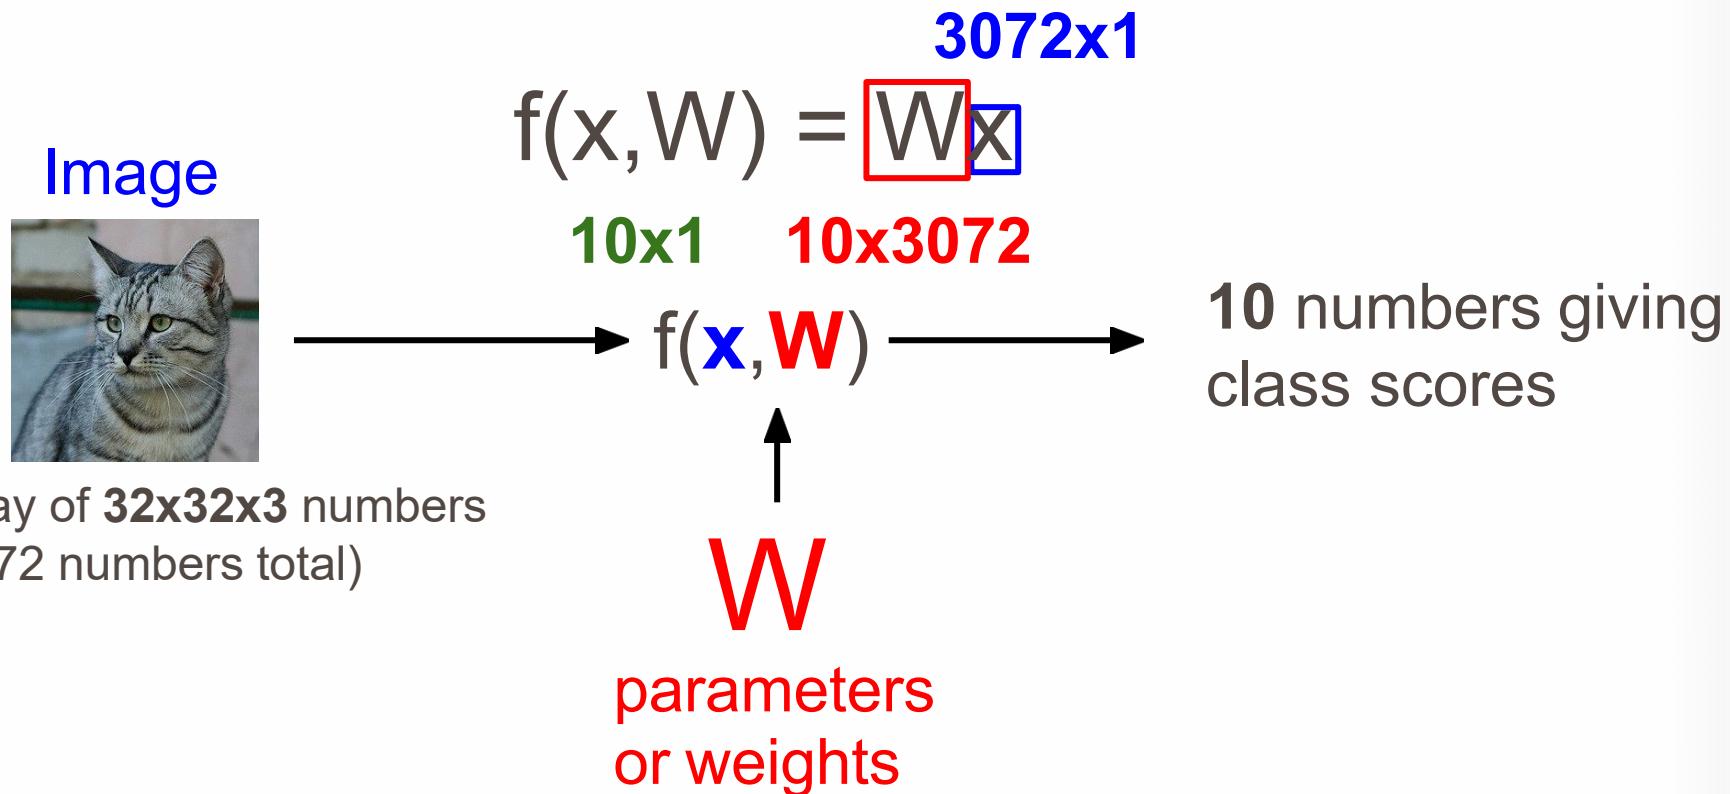
\mathbf{W}

parameters
or weights

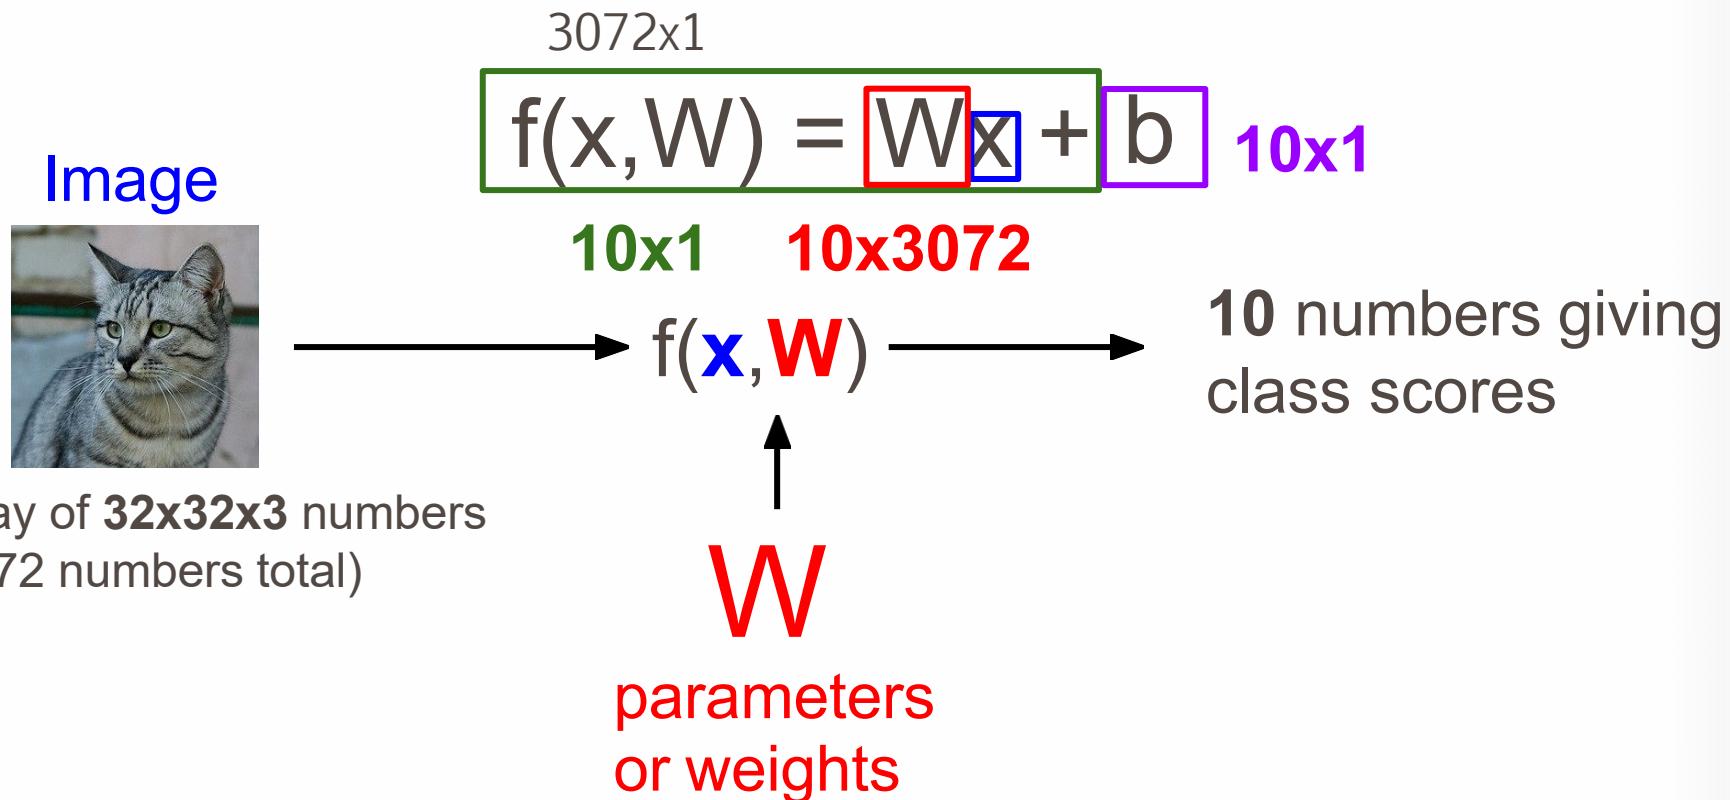
$$f(\mathbf{x}, \mathbf{W})$$

10 numbers giving
class scores

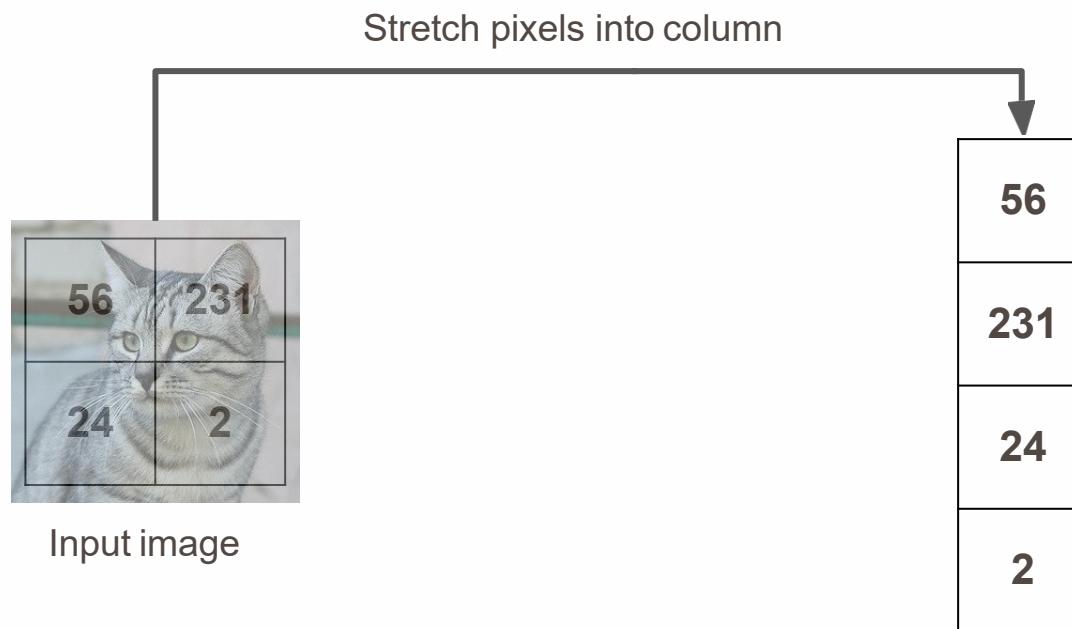
Parametric Approach: Linear Classifier



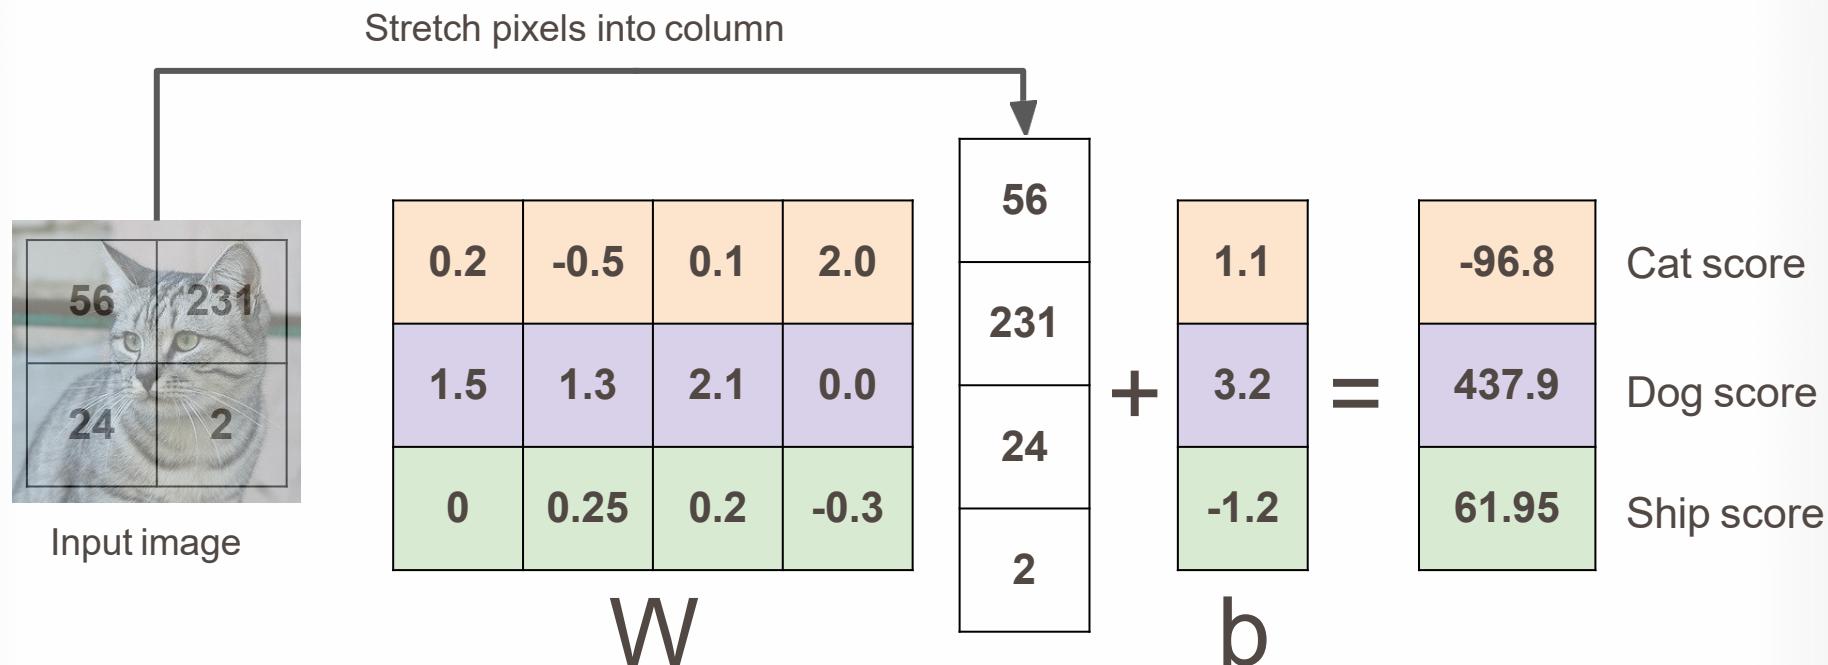
Parametric Approach: Linear Classifier



Example with an image with 4 pixels, and 3 classes
(cat/dog/ship)



Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



Example with an image with 4 pixels, and 3 classes (cat/dog/ship)

Algebraic Viewpoint

$$f(x, W) = Wx$$

Stretch pixels into column

$$\begin{array}{c}
 \text{Input image} \\
 \begin{array}{|c|c|} \hline
 56 & 231 \\ \hline
 24 & 2 \\ \hline
 \end{array}
 \end{array}
 \xrightarrow{\text{Stretch pixels into column}}
 \begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 56 \\ \hline
 231 \\ \hline
 24 \\ \hline
 2 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|c|c|} \hline
 0.2 & -0.5 & 0.1 & 2.0 \\ \hline
 1.5 & 1.3 & 2.1 & 0.0 \\ \hline
 0 & 0.25 & 0.2 & -0.3 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 0.2 & -0.5 \\ \hline
 0.1 & 2.0 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 1.5 & 1.3 \\ \hline
 2.1 & 0.0 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 0 & 0.25 \\ \hline
 0.2 & -0.3 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|} \hline
 1.1 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 -96.8 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 437.9 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 61.95 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 0.2 & -0.5 \\ \hline
 0.1 & 2.0 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 1.5 & 1.3 \\ \hline
 2.1 & 0.0 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 W \\
 \begin{array}{|c|c|} \hline
 0 & 0.25 \\ \hline
 0.2 & -0.3 \\ \hline
 \end{array}
 \end{array}$$

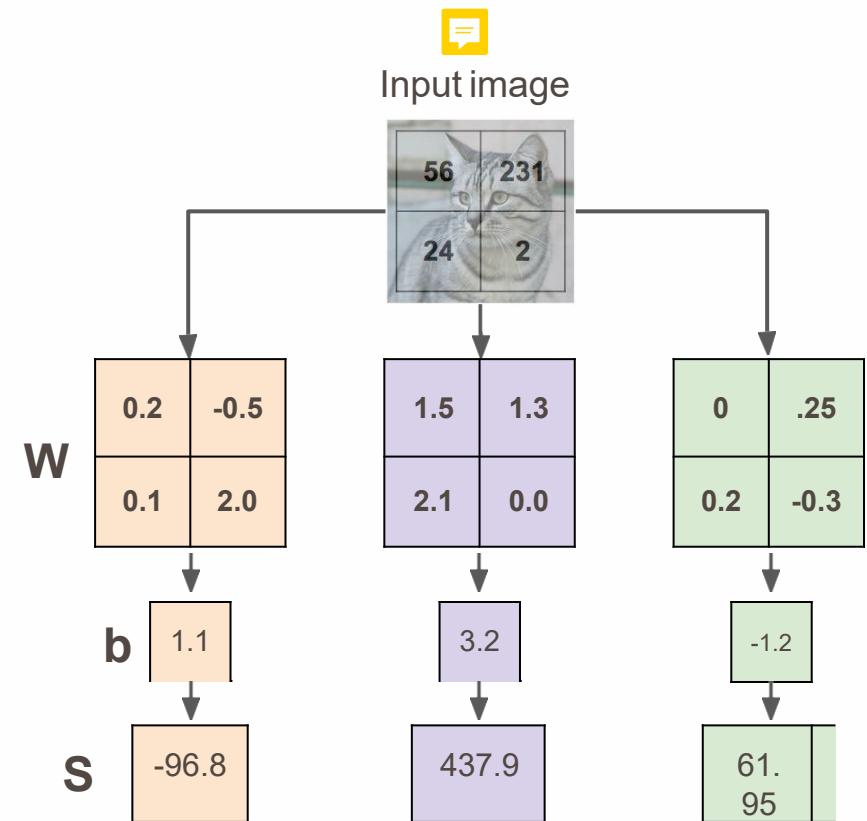
$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 3.2 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 -1.2 \\ \hline
 \end{array}
 \end{array}$$

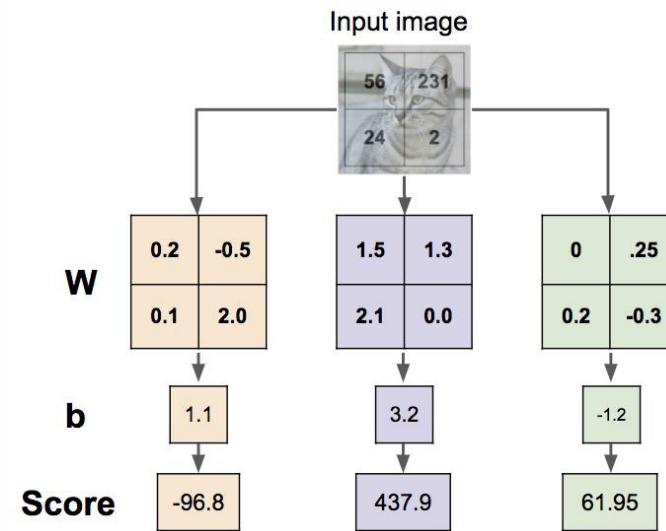
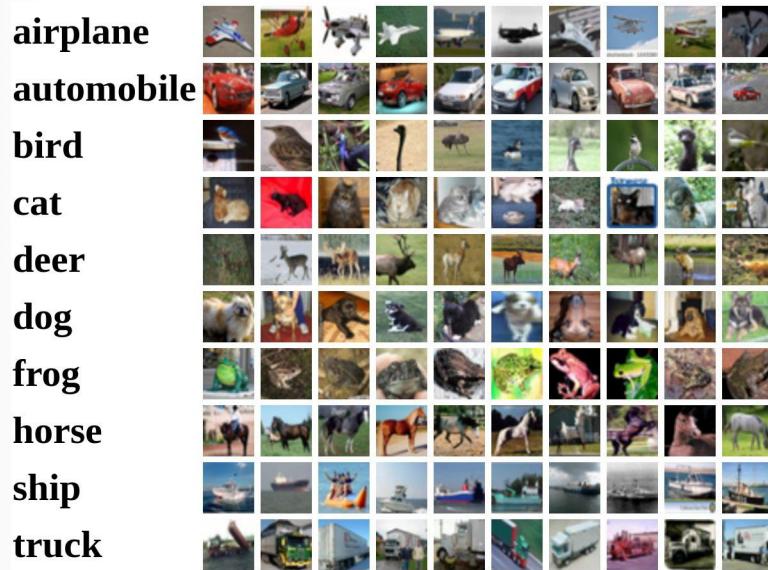
$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 -96.8 \\ \hline
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 437.9 \\ \hline
 \end{array}
 \end{array}$$

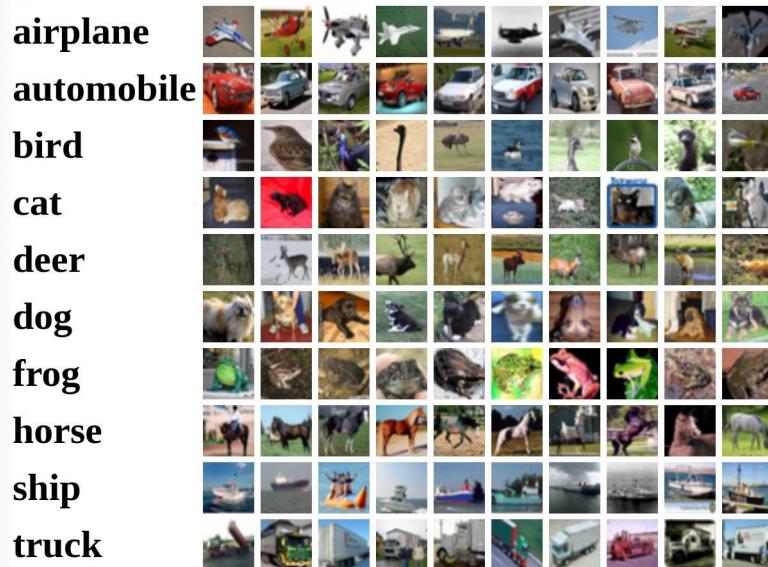
$$\begin{array}{c}
 b \\
 \begin{array}{|c|} \hline
 61.95 \\ \hline
 \end{array}
 \end{array}$$



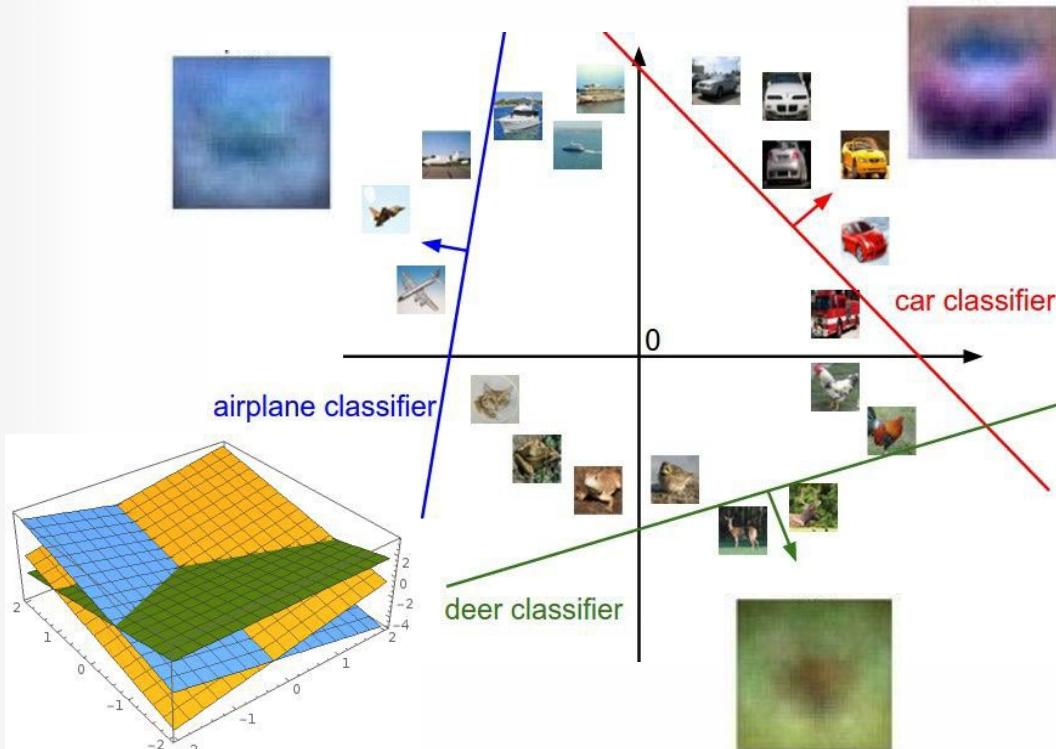
Interpreting a Linear Classifier



Interpreting a Linear Classifier: Visual Viewpoint



Interpreting a Linear Classifier: Geometric Viewpoint



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers
(3072 numbers total)

Plot created using [Wolfram Cloud](#)

Cat image by [Nikita](#) is licensed under [CC-BY 2.0](#)

Hard cases for a linear classifier

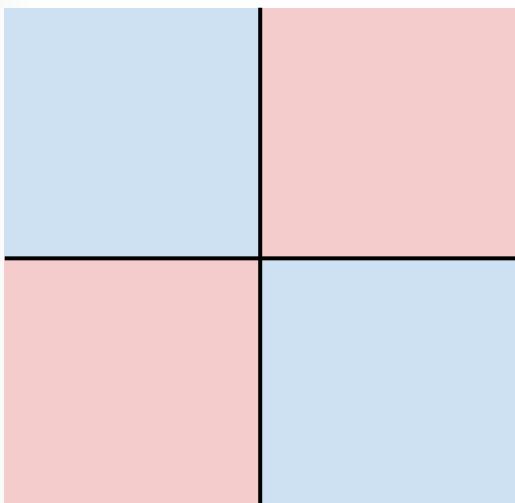


Class 1:

First and third quadrants

Class 2:

Second and fourth quadrants

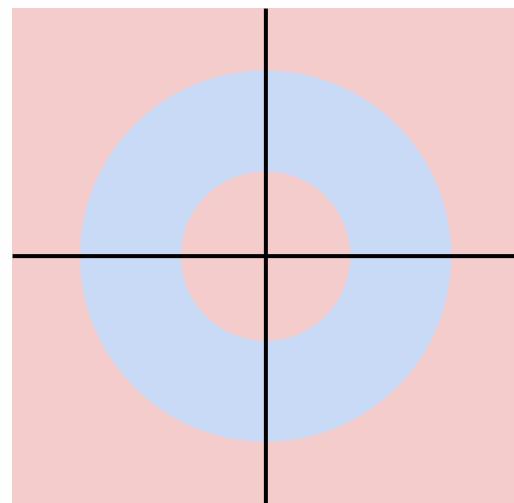


Class 1:

$1 \leq \text{L2 norm} \leq 2$

Class 2:

Everything else

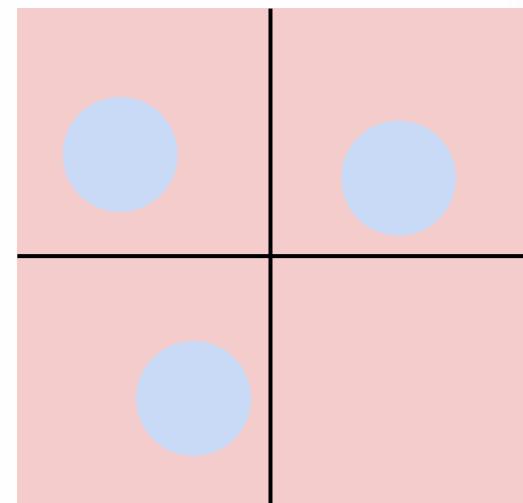


Class 1:

Three modes

Class 2:

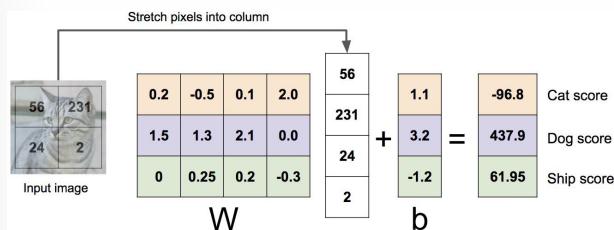
Everything else



Linear Classifier: Three Viewpoints

Algebraic Viewpoint

$$f(x, W) = Wx$$



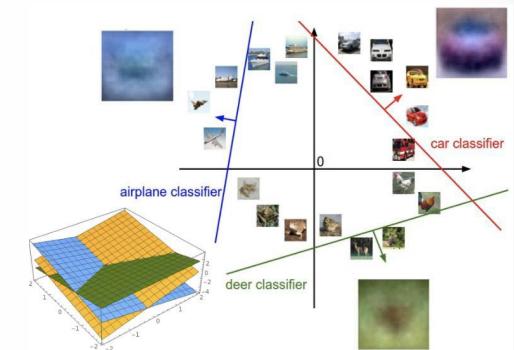
Visual Viewpoint

One template per class



Geometric Viewpoint

Hyperplanes cutting up space



Linear Classifier?

Defined a (linear) score function $f(x, W) = Wx + b$

Example class
scores for 3
images for some W

How can we tell
whether this W is
good or bad?



airplane	-3.45	-0.51	3.42
automobile	-8.87	6.04	4.64
bird	0.09	5.31	2.65
cat	2.9	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	-4.34
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

Cat image by [Nikita](#) is licensed under [CC-BY 2.0](#)

Car image is [CC0 1.0](#) public domain

Frog image is in the public domain

Linear Classifier: $f(x, W) = Wx + b$

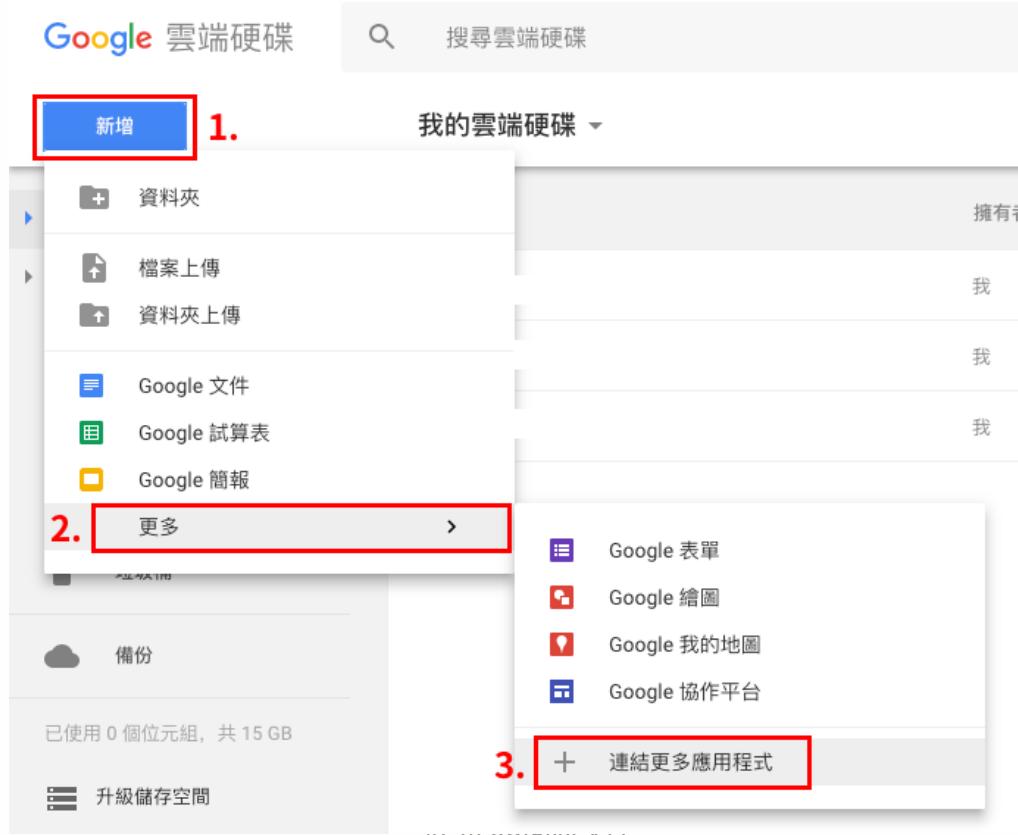
- Next:
 - Loss function
 - quantifying what it means to have a good W
 - Optimization
 - start with random W and find a W that minimizes the loss
 - CNNs!
 - tweak the functional form off



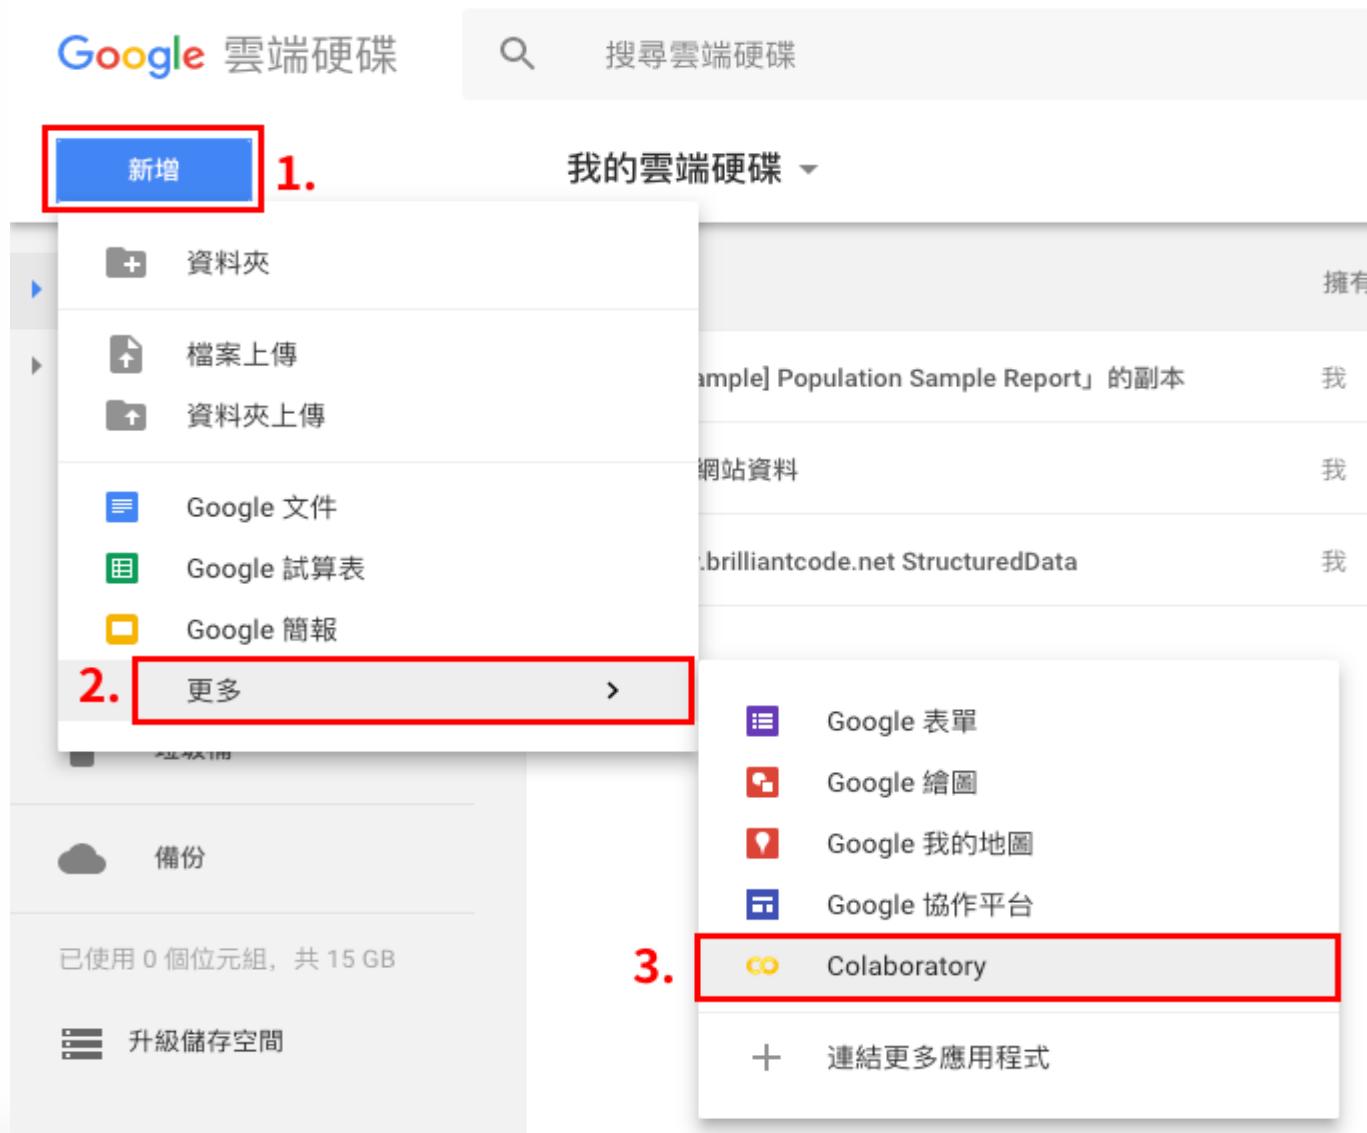
GOOGLE CLOUD RESOURCE: COLAB

Run TF without GPU?

- Google COLAB (Colaboratory)
 - Only three steps
 - Having free GPU/TPU resource if you have a google account



Click right key on ipynb and find colab

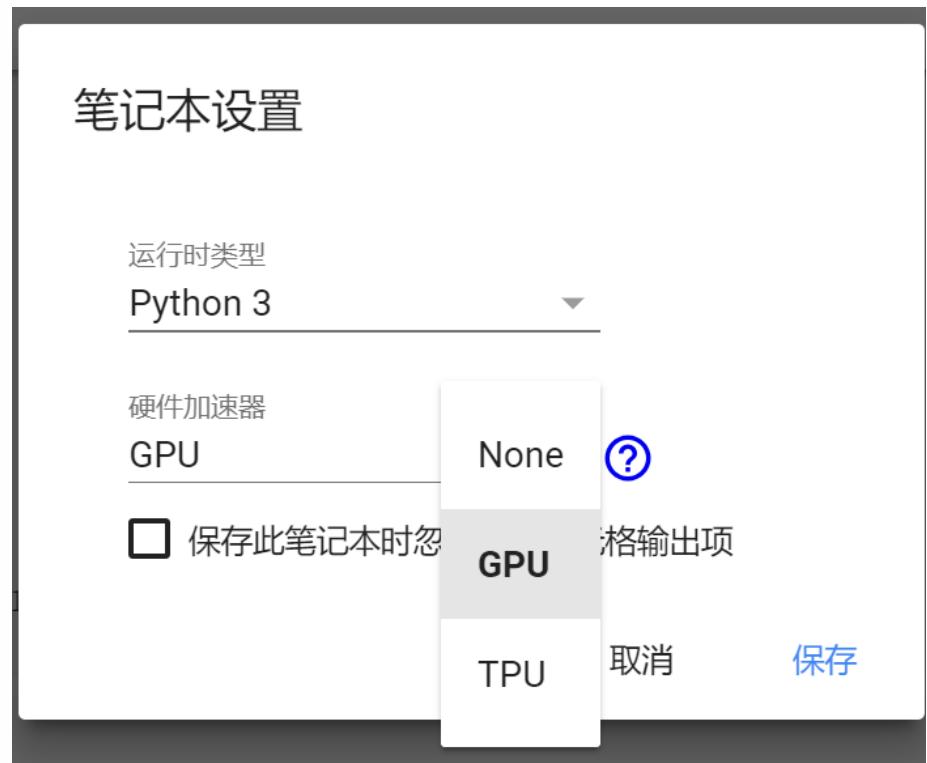
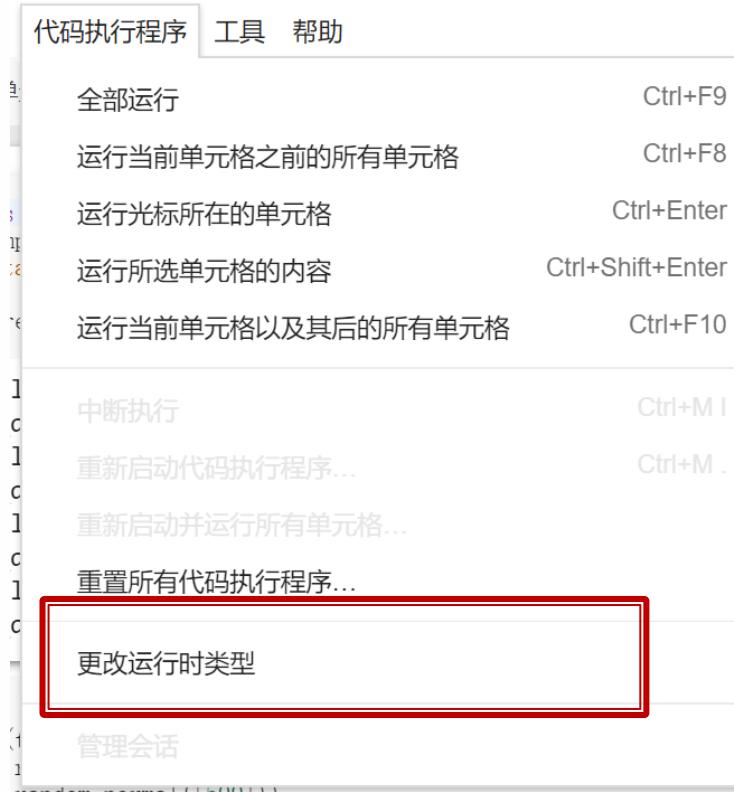


COLAB

- Free GPU (K80/V100) and TPU resource
 - 24G RAM (very large, compared to standard GPU)
 - RTX 2080 Ti
 - NTD 35,000
 - K80
 - NTD 180, 000
- You only can run a program on the free GPU 12 hours per day
 - Wait another 12 hours to access the free resource
- Runtime environment
 - Can be CPU/GPU/TPU

GPU Resource Usage

ipynb ☆

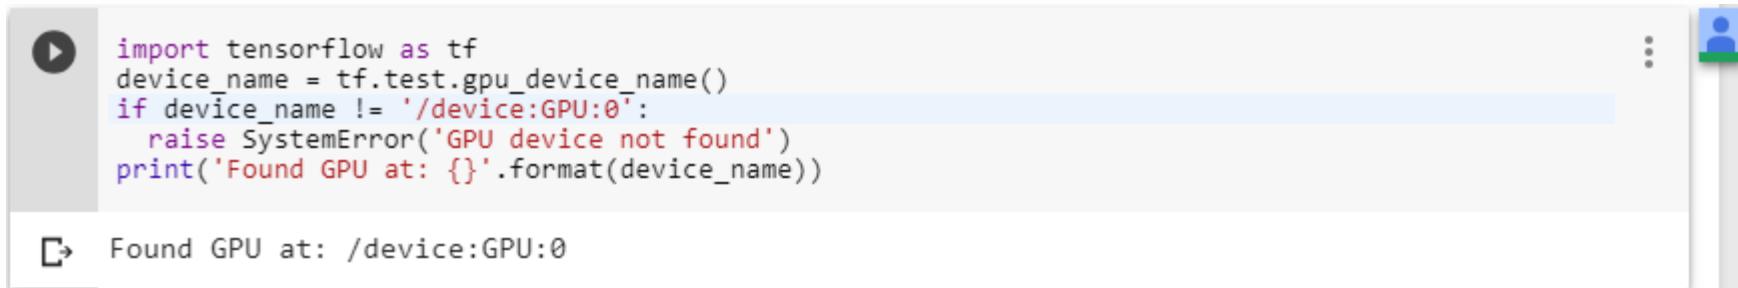


GPU vs TPU? TPU 為 Google 自家提出的加速硬體，照道理應該會比較快！大家亦可試試看。

Check Whether The GPU Resource is Enabled

```
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))
```

- 若有使用到GPU，則會顯示



The screenshot shows a Jupyter Notebook cell with the following content:

```
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))
```

Below the code cell, the output is displayed:

Found GPU at: /device:GPU:0

COLAB with Your Own Files

```
from google.colab import drive  
import os  
  
drive.mount('/content/gdrive' )  
os.chdir("/content/gdrive/My Drive") #更改路徑  
os.getcwd() #查看當前路徑
```

Use ! To run bash shell in colab

```
!ls  
## show your files
```

- 將上述程式碼貼到 COLAB 中，並根據【畫面指示】操作，就可以有與 GDRIVE 連線的權限

Connect Google Drive with COLAB

```
!mkdir -p Drive  
!google-drive-ocamlfuse Drive
```

- 上述指令在於
 - 在 COLAB 中建立一個資料夾 Drive (遠端中)
 - 利用指令!google-drive-ocamlfuse，把你連結到的雲端硬碟，連結到 Drive 資料夾中
 - 換句話說，你存取 Drive 資料夾，等於存取你的雲端硬碟

```
!ls Drive
```

```
00125870.pdf  
00125870.pdf.odt  
'00125870.pdf 的翻譯版本.odt'  
'1030909-清大(電機工程學系 ) SERVER.xlsx'  
'1030909-清大(電機工程學系 ) SERVER.xlsx.ods'  
'1030910-清大(電機工程學系 ) SERVER.pdf'
```



GOOGLE CLOUD SERVICE: VISION (AUTOML)

線上免費模型 (12個月免費)

- <https://cloud.google.com/vision>



AI & Machine Learning Products

Cloud Vision

運用強大且經過預先訓練的 API 模型，從圖片中獲取深入分析資料，也可以使用 AutoML Vision^{測試版}輕鬆訓練自訂的視覺模型。

 免費試用

[查看這項產品的說明文件。](#)

Cloud Vision

運用強大且經過預先訓練的 API 模型，從圖片中獲取深入分析資料，也可以使用 AutoML Vision^{測試版}輕鬆訓練自訂的視覺模型。

[前往主控台](#)

[查看這項產品的說明文件。](#)

AutoML 授權申請 (尚未有付費帳號)

免費試用 Google Cloud Platform

步驟 2 之 1

國家/地區

台灣

服務條款

- 我同意《[Google Cloud Platform 服務條款](#)》和[所有適用服務和 API](#)的服務條款。我也已詳閱並同意遵守《[Google Cloud Platform 免費試用版的服務條款](#)》。

必須勾選這個核取方塊才能提交表單

電子郵件最新消息

- 我想要定期收到 Google Cloud 和 Google Cloud Partner 傳送的電子郵件，隨時掌握相關新聞、產品動態和特價優惠。

同意並繼續



AutoML 授權申請 (尚未有付費帳號)

- 必須有信用卡
 - 不會偷偷被扣款
- 一般銀行的卡?
 - 台灣尚不支援
- 其他方法?
 - 請自行Google

付款方式



每月自動付款

您採用的付費方式是每月結帳日定期自動扣款。

付款方式 ⓘ

卡號

#

卡號為必填

持卡人姓名

許志仲

月份 / 年份

CVC

信用卡或簽帳金融卡地址同上

Cloud Vision (AutoML) 使用方法

1 Upload and label images



2 Train your model



3 Evaluate



■ 首先建立資料庫 (資料夾: 類別名稱)

名稱	修改日期	類型
5498592	2019/3/15 上午 0...	檔案資料夾
5513628	2019/3/15 上午 0...	檔案資料夾
5540978	2019/3/15 上午 0...	檔案資料夾

AutoML訓練自己資料庫方法

- 建立 Cloud Storage Bucket (CSB)
- 建立 CSB 的權限
- 上傳資料到自己的CSB
 - Including image files and their labels (.csv)
- 建立自己 CSB 的 Training dataset
- 訓練模型 & Testing

AutoML: 建立 CSB

- 首先開啟 Cloud AutoML 以及 Storage 的 API

[http://console.cloud.google.com/?cloud](http://console.cloud.google.com/?cloudshell=true)

- 首先連結到 ^{shell=true} Google cloud 主控台 Console

- PROJECT=\$(gcloud config get-value project) && BUCKET="\${PROJECT}-vcm"
- gsutil mb -p \${PROJECT} -c regional -l us-central1 gs://\${BUCKET}

- 建立 Service

```
PROJECT=$(gcloud config get-value project)
```

```
gcloud projects add-iam-policy-binding $PROJECT \  
  member="serviceAccount:custom-vision@appspot.gserviceaccount.com" \  
  role="roles/ml.admin"
```

```
gcloud projects add-iam-policy-binding $PROJECT \  
  member="serviceAccount:custom-vision@appspot.gserviceaccount.com" \  
  role="roles/storage.admin"
```

建立自己的資料庫

■ 進入到 AutoML Vision 的頁面

AutoML Vision **BETA**

faces ADD IMAGES LABEL STATS EXPORT DATA

IMAGES TRAIN EVALUATE PREDICT

	All images	66
Q	Labeled	66
Q	Unlabeled	0
	Type to filter...	⋮
	5498592	22
	5513628	16
	5540978	28
	Add label	

Label: 5513628 X

Select all images

Type to filter...

 5513628  5513628  5513628  5513628

 5513628  5513628  5513628  5513628

訓練MODEL

AutoML Vision **BETA** faces **ADD IMAGES** **LABEL STATS** **E**

IMAGES **TRAIN** EVALUATE PREDICT

Try labeling more images before training

Each label should have at least 100 images assigned. Fewer images often result in inaccurate precision and recall scores. [Learn more](#)

5498592	22
5513628	16
5540978	28

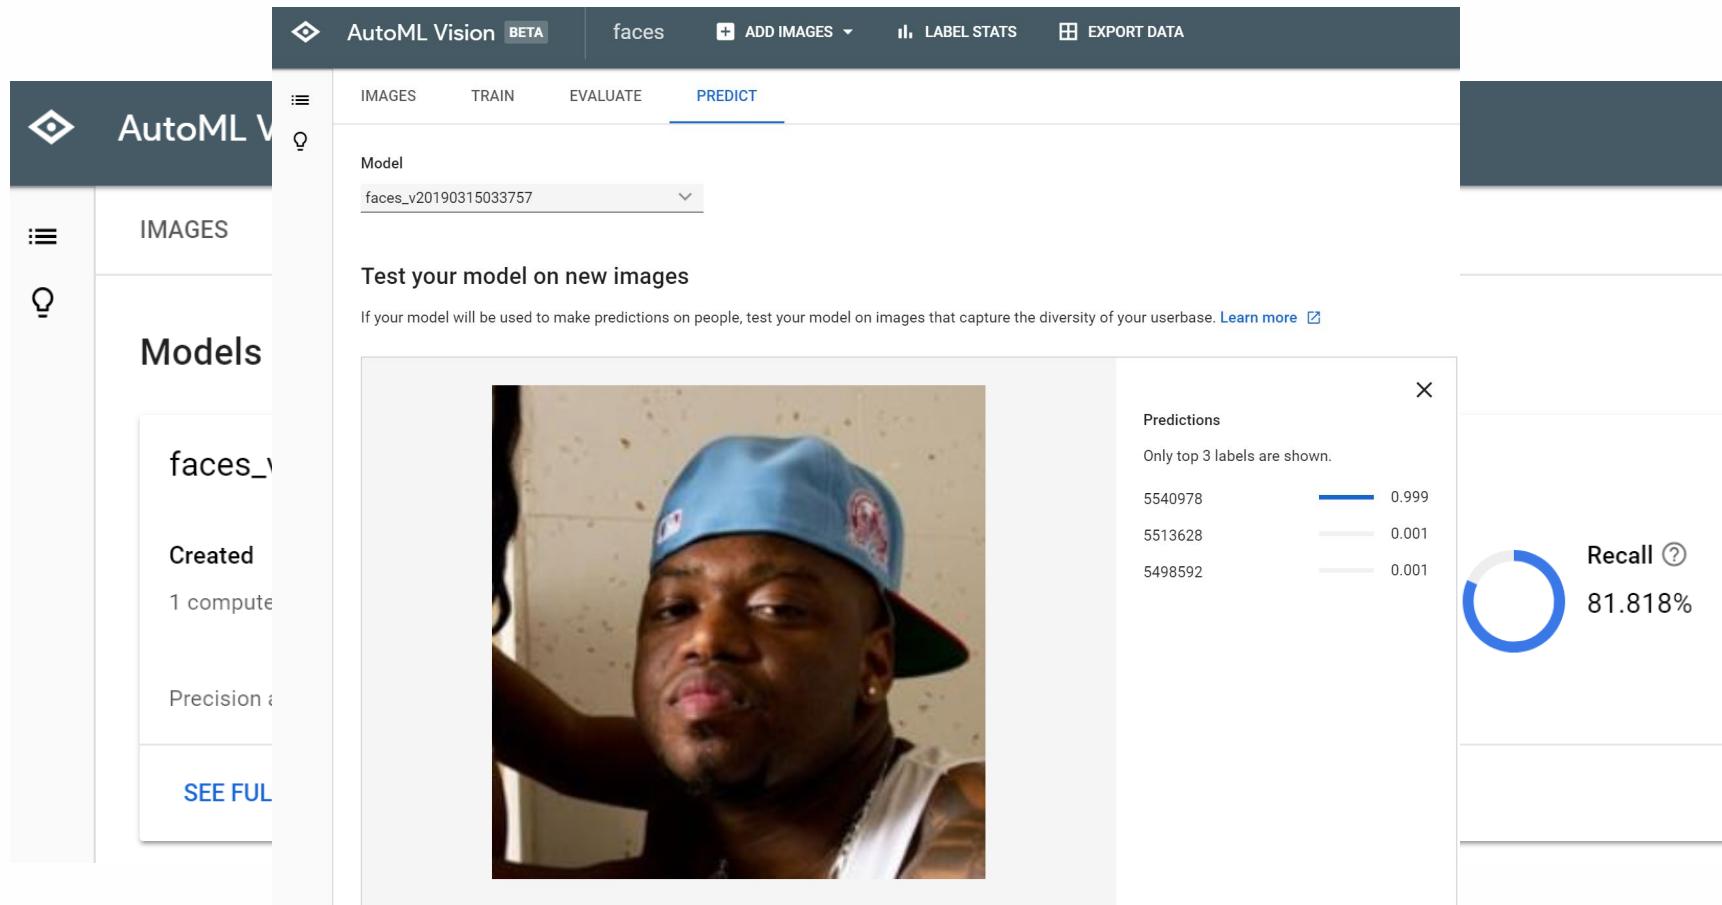
Your images will be automatically split into [training and test sets](#), so you can evaluate your model's performance. Unlabeled images will not be used.

Training images	46
Validation images	9
Test images	11

START TRAINING

Learned Model

- 可以馬上看Training 結果，Predict 可以上傳照片並辨識資料



Comparison

	Tensorflow (Own PC)	COLAB (Tensorflow)	AutoML
執行速度	快	中	中 (要網路)
訓練速度	快	中	快 (已上傳的話)
建立方法速度	慢	慢	超快
方便性	一般 (要寫程式)	優 (有網路就可)	優 (要註冊，未來要錢)
客製化彈性	優 (自己建立模型)	優	無
辨識效能	高	高	高一點
任務置換	容易	容易	麻煩 (要換資料)
I/O效能	高 (SSD版)	低 (有限制)	很高
成本	一般 (GPU的錢)	低 (只要網路)	中 (算訓練時間收費)